

Introduction to Graph Algorithms

Data Structures and Algorithms (094224)

Yuval Emek

Winter 2022/23

The role of graphs

- A fundamental combinatorial structure
- A mathematical model for communication networks, street and road maps, power grids, VLSI circuits, biological cellular networks
- Abstraction for **pairwise relations** between the components of a system
- Fundamental to discrete algorithms
 - An abundance of algorithms on (and data structures for) graphs
 - Provide useful insights for many other algorithms and data structures

1 Graph theory — basic definitions

- Trees



2 Graph representation

3 Eulerian cycles

Directed and Undirected graphs

- **Graph** $\langle \text{גרף} \rangle G = (V, E)$
 - V = a set of *vertices* (צמתים) (or *nodes* (קודקודים))
 - E = a set of vertex pairs referred to as *edges* (קשתות) (or *arcs* (צלעות))
 - Unless stated otherwise, our graphs are *finite* (סופיים): V, E finite sets
- **Directed** graph (or *digraph*):
 - The edges are ordered pairs
 - $(u, v) \neq (v, u)$
 - $E \subseteq V \times V$ is a (binary) relation over V
 - Edge (v, v) is possible — called a *self-loop* (לולאה עצמית)
 - The directed graph is called *simple* (פשוט) if it doesn't have self-loops
- **Undirected** graph:
 - The edges are unordered pairs
 - $(u, v) = (v, u)$
 - Similar to symmetric relation, but strictly speaking, not a relation at all
 - Unless stated otherwise, no self-loops
- (Un)directed graph: u, v are the *endpoints* (קצנות) of edge (u, v)

Pictorial representation of graphs



Incidence, adjacency, and degrees

- $e = (u, v) \in E$ in a directed graph:
 - e is *incident from* or *leaves* \langle ויצאתה \rangle u and *incident to* or *enters* \langle יכונסת \rangle v
 - v is *adjacent* \langle סמוך \rangle to u
- $e = (u, v) \in E$ in an undirected graph:
 - e is *incident on* \langle נוןעtha \rangle u and v
 - u and v are *adjacent* \langle סמכים \rangle to each other
 - u and v are *neighbors* \langle שכנים \rangle
- Directed graph:
 - *out-degree* \langle דרגת יציאה \rangle of v = #edges incident from v
 - *in-degree* \langle דרגת כניסה \rangle of v = #edges incident to v
 - *degree* \langle דרגה \rangle of v = in-degree of v + out-degree of v
- Undirected graph:
 - *degree* \langle דרגה \rangle of v = #edges incident on v
- Questions:
 - In directed graphs, what is $\sum_{v \in V} \deg_{\text{in}}(v)$? $\sum_{v \in V} \deg_{\text{out}}(v)$?
 - In undirected graphs, what is $\sum_{v \in V} \deg(v)$? $= 2 * \text{number of edges}$

Paths

- The vertex sequence $P = \langle v_0, v_1, \dots, v_k \rangle$ is a **path** **מסלול** in $G = (V, E)$ if $(v_{i-1}, v_i) \in E$ for every $1 \leq i \leq k$
 - The path leads **from** v_0 **to** v_k
- If such a path P exists, then v_k is **reachable** **נגיש** from v_0
- The **length** **אורך** of path P is k
- We say that path P **contains** **מכיל** vertices v_0, v_1, \dots, v_k and edges $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$
- Path P is **simple** **פשוט** if v_0, v_1, \dots, v_k are distinct
- A **subpath** **חת-מסלול** of path P is a contiguous subsequence of the vertex sequence $\langle v_0, v_1, \dots, v_k \rangle$
 - A path by its own right

From a path to a simple path

Lemma

If $G = (V, E)$ admits a path from u to v , then $G = (V, E)$ admits a simple path from u to v .

Proof. Let $P = \langle x_0, x_1, \dots, x_k \rangle$ be a **shortest path** from $u = x_0$ to $v = x_k$

- Assume by **contradiction** that P is not simple
 - $x_i = x_j$ for some $0 \leq i < j \leq k$
- Let $P' = \langle x_0, x_1, \dots, x_i = x_j, x_{j+1}, \dots, x_k \rangle$
- P' is also a path from u to v
- $\text{length}(P') < \text{length}(P)$
- Contradiction! \implies assumption must be wrong ■

- Path $C = \langle v_0, v_1, \dots, v_k \rangle$ in a directed graph forms a **cycle** $\langle \text{מעגל} \rangle$ if
 - $v_0 = v_k$; and
 - $k \geq 1$
- Path $C = \langle v_0, v_1, \dots, v_k \rangle$ in an undirected graph forms a **cycle** $\langle \text{מעגל} \rangle$ if
 - $v_0 = v_k$;
 - $k \geq 3$; and
 - $v_{i-1} \neq v_{i+1}$ for every $0 < i < k$
- Directed graph: cycle contains at least 1 edge
Undirected graph: cycle contains at least 3 edges
- Cycle C is **simple** $\langle \text{פשוט} \rangle$ if v_1, \dots, v_k are distinct
- Paths $\langle v_0, v_1, \dots, v_{k-1}, v_0 \rangle$ and $\langle v'_0, v'_1, \dots, v'_{k-1}, v'_0 \rangle$ form the same cycle if $\exists j \in \mathbb{Z}$ such that $v'_i = v_{i+j} \pmod{k}$ for $i = 0, 1, \dots, k-1$
- An (un)directed graph with no cycles is called **acyclic** $\langle \text{חסר מעגליים} \rangle$

Subgraphs

- (Un)directed graph $G = (V, E)$
- $G' = (V', E')$ is a *subgraph* חתגראף of G if $V' \subseteq V$ and $E' \subseteq E$
- *graph by its own right*
- The subgraph of G *induced* מושרה by $V' \subseteq V$ is $G' = (V', E')$, where $E' = \{(u, v) \in E \mid u, v \in V'\}$
- A subgraph $G' = (V', E')$ of G is *spanning* פורש if $V' = V$

Connectivity

- Undirected graph $G = (V, E)$
 - Called *connected* $\langle \text{קשור} \rangle$ if v is reachable from u for every $u, v \in V$
 - The equivalence classes of the “is reachable from” relation on the vertices are called *connected components* $\langle \text{רכיבים קשורים} \rangle$
- Directed graph $G = (V, E)$
 - Called *strongly connected* $\langle \text{קשור דיטב} \rangle$ if v is reachable from u for every $u, v \in V$
 - The equivalence classes of the “are mutually reachable” relation on the vertices are called *strongly connected components* $\langle \text{רכיבים קשורים דיטב} \rangle$

Number of edges in connected graphs

Lemma

If undirected graph $G = (V, E)$ is connected, then $|E| \geq |V| - 1$.

Proof. By induction on $|V|$.

- Base: holds (with equation) for $|V| = 1$
- Hypothesis: holds for any graph with n vertices
- Step: consider a connected graph $G = (V, E)$ with $|V| = n + 1$
 - Assume by contradiction that $|E| < |V| - 1$
 - Sum of degrees $< 2|V| - 2 \implies$ average degree < 2
 - Pigeonhole principle: $\exists v \in V$ of degree $\deg(v) < 2$
 - $\deg(v) = 1$ (can't be 0 as G is connected)
 - $G' =$ the subgraph of G induced by $V \setminus \{v\}$

The proof continues

- Inductive step — cont.
 - G' is **connected**
 - $\deg(v) = 1$ means that v can't be an **interior vertex** in any simple path
 - Removing it from G doesn't affect reachability between other vertices
 - G' has $|V| - 1$ vertices and $|E| - 1$ edges
 - **Inductive hypothesis:** $|E| - 1 \geq |V| - 2 \implies |E| \geq |V| - 1$
 - Contradiction! \implies assumption must be wrong ■

(Un)directed versions

- The *directed version* \langle גרסת מכוונת \rangle of an undirected graph $G = (V, E)$ is the directed graph $G' = (V', E')$ such that
 - $V' = V$
 - $E' = \{(u, v) \mid (u, v) \in E\}$
 - Observe: $|E'| = 2|E|$
- The *undirected version* \langle גרסת לא מכוונת \rangle of a directed graph $G = (V, E)$ is the undirected graph $G' = (V', E')$ such that
 - $V' = V$
 - $E' = \{(u, v) \mid (u, v) \in E, u \neq v\}$
 - Observe: $|E'| \leq |E|$

Important graph classes

Undirected graphs:

- A *complete* $\langle \text{מלא} \rangle$ graph $G = (V, E)$:
 - $E = \{(u, v) \mid u, v \in V, u \neq v\}$
- A *bipartite* $\langle \text{דו צדדי} \rangle$ graph $G = (V, E)$:
 - V can be partitioned into $V = V_1 \cup V_2$, $V_1 \cap V_2 = \emptyset$, such that $E \subseteq V_1 \times V_2$
- An acyclic undirected graph is called a *forest* $\langle \text{יער} \rangle$
- A connected forest is called a *tree* $\langle \text{עץ} \rangle$

Directed graphs:

- A directed acyclic graph is abbreviated (and called) *DAG*

1 Graph theory — basic definitions

- Trees

2 Graph representation

3 Eulerian cycles

Characterizing trees

Theorem

The following statements are equivalent for undirected graph $G = (V, E)$:

- ① G is a tree
- ② Any two vertices in G are connected by a unique simple path
- ③ G is connected, but if any edge is removed, it disconnects
- ④ G is connected and $|E| = |V| - 1$
- ⑤ G is acyclic and $|E| = |V| - 1$
- ⑥ G is acyclic, but if any edge is added, a cycle is formed

Proof: by showing that $(1) \Rightarrow (2) \Rightarrow (3) \Rightarrow (4) \Rightarrow (5) \Rightarrow (6) \Rightarrow (1)$

Tree characterization proof (1) \Rightarrow (2)

- By (1): G is connected
 - \implies any two vertices in V admit ≥ 1 simple paths
- Left to show: inequality is not strict
- Assume by contradiction: G admits simple (u, v) -paths $p_1 \neq p_2$
- Let x be the vertex at which p_1 and p_2 first diverge
 - Exists because u appears in both p_1 and p_2 and $p_1 \neq p_2$
- Let y be the vertex following x at which p_1 and p_2 reconverge
 - Exists because v appears in both p_1 and p_2
- Let p'_i be the (x, y) -subpath of p_i for $i = 1, 2$
- p'_1 and p'_2 share no vertices other than their endpoints
- \implies path obtained by concatenating p'_1 and reverse of p'_2 is a cycle
- But G is a tree due to (1), hence cycle free ($\rightarrow \leftarrow$)

Tree characterization proof (2) \Rightarrow (3)

- By (2): any two vertices in V are connected by (unique) simple path
 - $\Rightarrow G$ is connected
- Left to show: removal of edge $e = (u, v) \in E$ disconnects G
- (u, v) is a path between u and v in G
 - \Rightarrow must be the unique simple (u, v) -path due to (2)
- Thus, if edge e is removed, then no simple (u, v) -path in G
 - \Rightarrow no (u, v) -path in G
- Therefore, removal of e disconnects G

Tree characterization proof (3) \Rightarrow (4)

- By (3): G is connected
 - $\Rightarrow |E| \geq |V| - 1$
- Prove that $|E| \leq |V| - 1$ by **induction** on $|V|$
- **Base:** connected graph with $n = 1$ or $n = 2$ vertices has $n - 1$ edges
- Let $|V| = n \geq 3$
- **Assume:** all graphs satisfying (3) with $< n$ vertices also satisfy $|E| \leq |V| - 1$
- Construct G' by **removing** an arbitrary edge $e \in E$ from G
 - Exists because $|E| \geq n - 1 \geq 2$
- G' has $k \geq 2$ **connected components** due to (3)
 - (Actually, $k = 2$)
- Each connected component $G_i = (V_i, E_i)$ of G' satisfies (3)
 - Otherwise, G does not satisfy (3)
- **Inductive hypothesis:** $|E_i| \leq |V_i| - 1$
- Thus, $|E| - 1 = \sum_{i=1}^k |E_i| \leq \sum_{i=1}^k (|V_i| - 1) = n - k$
 - $\Rightarrow |E| \leq n - k + 1 \leq n - 1$

Tree characterization proof (4) \Rightarrow (5)

- By (4): $|E| = |V| - 1$
- Left to show: G is acyclic
- Assume by contradiction: G contains a cycle $\langle v_0, v_1, \dots, v_{k-1}, v_0 \rangle$
 - W.l.o.g.: cycle is simple
- Let $G_k = (V_k, E_k)$ be the subgraph of G consisting of the cycle
 - $|V_k| = |E_k| = k$
- If $k < |V|$, then $\exists v_{k+1} \in V - V_k$ adjacent to some $v_i \in V_k$
 - G is connected by (4)
- Let $G_{k+1} = (V_{k+1}, E_{k+1})$ be the graph obtained from G_k by adding vertex v_{k+1} and edge (v_i, v_{k+1})
 - $|V_{k+1}| = |E_{k+1}| = k + 1$
- If $k + 1 < |V|$, then...
- ... obtain graph $G_n = (V_n, E_n)$
 - $V_n = V$
 - $|V_n| = |E_n| = n$
- $E_n \subseteq E \implies |E| \geq |E_n| = |V|$ ($\rightarrow \leftarrow$)

Tree characterization proof (5) \Rightarrow (6)

- By (5): G is acyclic
- Left to show: if any edge is added to G , then a cycle is formed
- Let k be #connected components of G
- Each connected component is a tree
 - Connected and acyclic
- Already proved (1) \Rightarrow (5), hence #edges in all connected components (together) is $|V| - k$
- By (5): $k = 1$, thus G is a tree
 - Connected and acyclic
- Already proved (1) \Rightarrow (2), hence any two vertices in V are connected by a unique simple path
- \implies adding any edge to G closes a cycle

Tree characterization proof (6) \Rightarrow (1)

- By (6): G is acyclic
- Left to show: G is connected
- Assume by contradiction: G has no path connecting vertices $u, v \in V$
- By (6): adding edge $e = (u, v)$ to E closes a cycle C
- Removing edge e from C yields (u, v) -path P
- By definition, P is included in $G_{(\rightarrow\leftarrow)}$ ■

Rooted trees

- A *rooted tree* (עץ מושרש) is a tree with one designated vertex
 - Called the *root* (שורש)
- Rooted tree $T = (V, E)$ with root r
- Any vertex on the unique simple path from v to r is called an *ancestor* (אב קדמון) of v
 - v is always an ancestor of v
- If u is an ancestor of v , then v is a *descendant* (צאצאי) of u
- The *subtree rooted at v* (חת-עץ מושרש) is the tree induced by the descendants of v rooted at v
- If u is the next vertex after v on the unique path from v to r , then u is the *parent* (אב) of v and v is a *child* (בן) of u
 - r is the only vertex in T with no parent
- A vertex with no children is a *leaf* (עליה)
- A non-leaf is an *internal vertex* (צומת פנימי)
- If two vertices share a parent, then they are *siblings* (אחים)

Rooted trees — cont.

- The *degree* (דֶּרֶגָה) of v in T = #children it has
 - Depends on whether we treat T as a graph or a rooted tree
- The *depth* (עומק) of v is the length of the unique path from v to r
 - What is the depth of r ?
- The *height* (גובה) of T is the maximum depth of any of its leaves
- The *height* (גובה) of v is the height of the subtree rooted at v
- A *level* (רמתה) of T consists of all vertices of the same depth
- The rooted tree is *ordered* (מוסדר) if the children of each internal vertex are ordered

וְיַדְעָה לִבְנֵי מִסְנָדָר

1 Graph theory — basic definitions

- Trees

2 Graph representation

3 Eulerian cycles

Representing graphs in a computer

- How do we represent (un)directed graph $G = (V, E)$ in a computer?
 - Prerequisite for graph algorithms
- Conventions:
 - $n = |V|$
 - Inside asymptotic notation, sometimes V
 - $m = |E|$
 - Inside asymptotic notation, sometimes E
 - In pseudocode, we write $G.V$ and $G.E$ (attributes of G)
- Convenient to assume that the vertices are numbered $1, \dots, n$
 - Arbitrary order
- Sometimes, vertices/edges have additional attributes (should be stored as well)
 - Common: edge *weights* $\langle w : E \rightarrow \mathbb{R} \rangle$
 - *Weighted graph* $\langle \text{נרכז מושקל} \rangle$

Adjacency matrix

- Matrix $A \in \{0, 1\}^{n \times n}$
 - $A(i, j) = \begin{cases} 1, & (i, j) \in E \\ 0, & \text{otherwise} \end{cases}$
 - Called *adjacency matrix* (מטריצה שכנות)
- Advantage: can answer the query “does $(i, j) \in E?$ ” in $\Theta(1)$ time
- If the graph is weighted, then $A \in \mathbb{R}^{n \times n}$
 - Depending on the context, $(i, j) \notin E$ represented by $A(i, j) = 0$, $A(i, j) = +\infty$, or $A(i, j) = -\infty$
- Entry $A(i, j)$ can include a pointer to an object with additional attributes of edge (i, j)
- How much memory does the adjacency matrix representation require?

Adjacency lists

- Array Adj of n lists
- For each $u \in V$, $\text{Adj}[u]$ is an *adjacency list* ⟨ שכניות ⟩ that contains an entry for each vertex $v \in V$ such that $(u, v) \in E$
 - The entry of v includes all attributes of edge (u, v)
 - Typically implemented with a (doubly) linked list
- Directed graph: edge (u, v) represented only in $\text{Adj}[u]$
Undirected graph: edge (u, v) represented in $\text{Adj}[u]$ and $\text{Adj}[v]$
 - In the directed case, can include a list for the entering edges too if necessary
- **Advantage:** can enumerate the edges incident on (resp., from) vertex v in $\Theta(\deg(v))$ (resp., $\Theta(\deg_{\text{out}}(v))$) time
 - How much time does it take in adjacency matrix representation?
- How much memory does the adjacency lists representation require?


1 Graph theory — basic definitions

- Trees

2 Graph representation

3 Eulerian cycles

The seven bridges of Königsberg



- The city of Königsberg (Prussia) included 7 bridges
- **The problem:** device a walk through the city that would cross every bridge exactly once
- Resolved negatively by **Leonhard Euler** in 1736
 - Considered as one of the origins of modern graph theory

The Eulerian cycle problem

- Undirected graph $G = (V, E)$
- A cycle C in G that goes through every edge in E exactly once is called an *Eulerian cycle* (מעגל אוילריאי)
- The Eulerian cycle problem:
 - Input: undirected graph $G = (V, E)$
Output: An Eulerian cycle in G or “no Eulerian cycle”
- An undirected graph that admits an Eulerian cycle is called an *Eulerian graph* (גרף אוילריאי)
 - If all degrees are positive, then G is connected

Characterizing Eulerian graphs

Theorem

A connected undirected graph is Eulerian iff the degree of every vertex is even.

- Connected undirected graph $G = (V, E)$
- For every cycle C in G and vertex $v \in V$,
 $\# \text{times } C \text{ enters } v = \# \text{times } C \text{ leaves } v$
- If C is Eulerian, then we can classify (all) edges incident on v as **entering** or **leaving** w.r.t. C
 - $\# \text{entering edges} = \# \text{leaving edges}$
- Impossible if $\deg(v)$ is odd
- Eulerian cycle \implies all degrees are even

Even degrees suffice

- Connected undirected graph $G = (V, E)$, all degrees are even
- Construct an Eulerian cycle C as follows:
 - ① Pick an arbitrary start vertex $v \in V$
 - ② Follow an arbitrary path P from v , erasing the edges from G upon traversing them, until getting back to v
 - Cannot get stuck at any vertex $\neq v$ because of the even degrees!
 - P forms a cycle
 - All degrees remain even!
 - ③ Merge P into C
 - ④ If C contains a vertex that includes incident edges that have not been removed from G , then call this vertex v and goto line 2
 - Finish after all edges are erased from G because G is connected
- All degrees are even \implies Eulerian cycle ■
- The proof is *constructive*
- Not a proper pseudocode
 - Missing data structure details
 - Missing details in line 3