

Efficient Data Structures

Data Structures and Algorithms (094224)

Yuval Emek

Winter 2022/23

Data structures revisited

- Focus: a dynamic set of objects
- Objects are typically organized in some **well structured digraph**
- A designated node attribute: **key** = the object's identifier
 - Keys are totally ordered
 - For simplicity, assume **uniqueness**
- Supporting (a subset of the) operations:
insert, delete, search, minimum, maximum, successor, predecessor
 - **FAST!!!**
- Actual data stored in satellite attributes
 - Sometimes a pointer to the actual data (storage considerations)

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
- Dynamic updates
- Discussion

3 Binary heaps

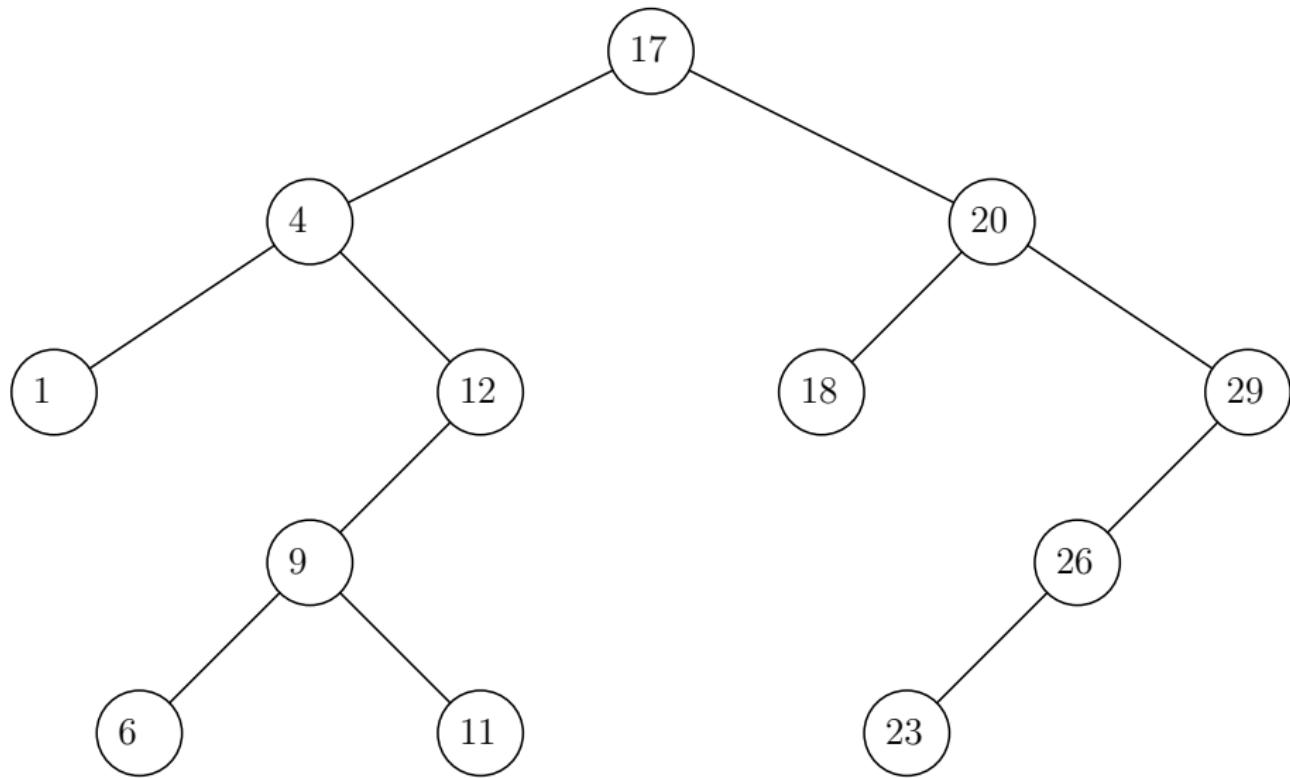
- Binary heap operations
- Discussion

A (rooted) binary tree organization

עץ חיפוש בינארי

- Objects associated with nodes of a rooted tree T with degrees ≤ 2
 - a.k.a. **binary tree**
- Data structure's attribute:
 - root = pointer to root of T
- Node attributes (additional):
 - left = pointer to left child in T
 - right = pointer to right child in T
 - p = pointer to parent in T
- The binary search tree property:
if y is a node in the left (resp., right) subtree of node x , then
 $y.\text{key} < x.\text{key}$ (resp., $y.\text{key} > x.\text{key}$)
 - Binary search tree is **not unique** for a given set of keys

Example



1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
- Dynamic updates
- Discussion

3 Binary heaps

- Binary heap operations
- Discussion

Searching

Search in the binary search tree rooted at x for a node whose key is k

`Tree_Search(x, k)`

```
1: if  $x == nil$  or  $x.key == k$  then
2:     return  $x$ 
3: if  $k < x.key$  then
4:     return Tree_Search( $x.left, k$ )
5: else
    return Tree_Search( $x.right, k$ )
```

- Run-time: $O(\text{height}(T))$

Finding the minimum/maximum element

Find the node with the smallest key in the binary search tree rooted at x

`Tree_Minimum(x)`

```
1: while  $x.\text{left} \neq \text{NIL}$  do
2:      $x = x.\text{left}$ 
3: return  $x$ 
```

- Undefined if the tree is empty
- Run-time: $O(\text{height}(T))$
- How do we find the maximum?

Finding a successor/predecessor

Find the node y with the smallest key among those with $y.key > x.key$

`Tree_Successor(x)`

```
1: if  $x.right \neq NIL$  then
2:     return Tree_Minimum( $x.right$ )
3:  $y = x.p$ 
4: while  $y \neq NIL$  and  $x == y.right$  do
5:      $x = y$ 
6:      $y = y.p$ 
7: return  $y$ 
```

- Run-time: $O(\text{height}(T))$
- How do we find a predecessor?

the opposite

Inserting a new node

Insert the new node z (DS attributes initialized to NIL) into T

`Tree_Insert(T, z)`

```
1: if  $T.\text{root} == \text{NIL}$  then
2:    $T.\text{root} = z$ 
3: else
4:    $y = T.\text{root}$ 
5:    $x = \text{NIL}$ 
6:   while  $y \neq \text{NIL}$  do
7:      $x = y$ 
8:     if  $z.\text{key} < y.\text{key}$  then
9:        $y = y.\text{left}$ 
10:    else  $y = y.\text{right}$ 
11:    $z.p = x$ 
12:   if  $z.\text{key} < x.\text{key}$  then
13:      $x.\text{left} = z$ 
14:   else  $x.\text{right} = z$ 
```

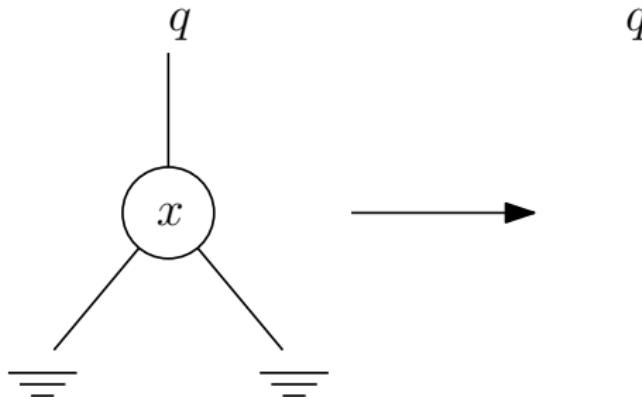
- Run-time: $O(\text{height}(T))$

Deleting a node

Deleting node x

case 1: x is a leaf

Not much to do



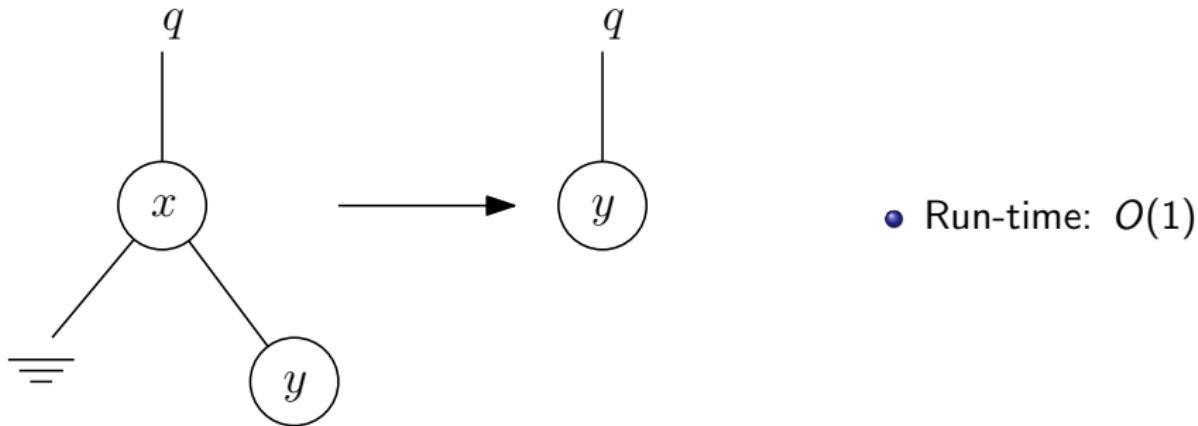
- Run-time: $O(1)$

Deleting a node — cont.

Deleting node x

case 2: x has a right child y and no left child

Replace x with y

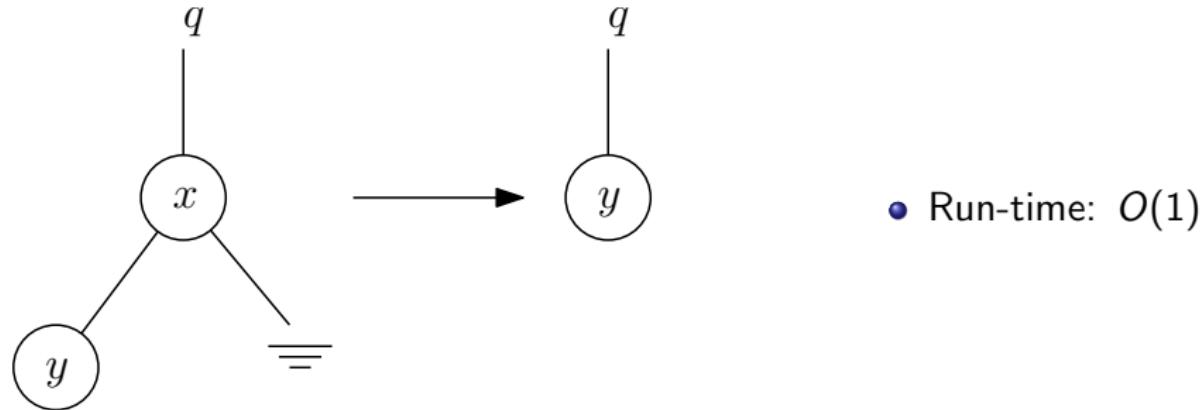


Deleting a node — cont.

Deleting node x

case 3: x has a left child y and no right child

Replace x with y

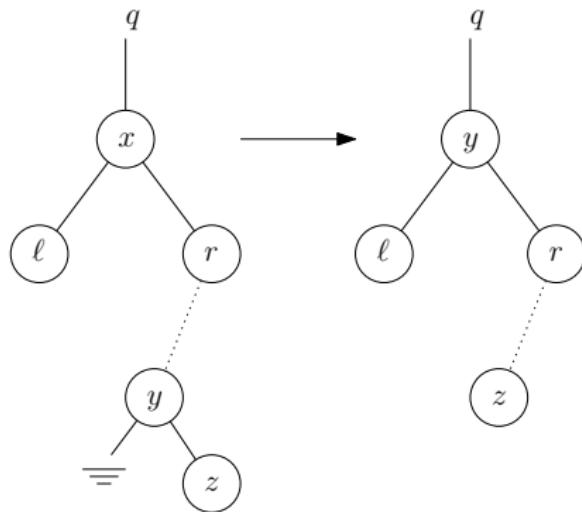


Deleting a node — cont.

Deleting node x

case 4: x has two children

- ① Find x 's successor y (in x 's right subtree)
- ② Swap x and y
- ③ Remove x
 - We know how to do it because now x has ≤ 1 child



• Run-time: $O(\text{height}(T))$

Balanced trees

- Supported operations:
insert, delete, search, minimum, maximum, successor, predecessor
- Run-time of each operation in a binary search tree T is $O(\text{height}(T))$
- $\text{height}(T)$ can be $\Omega(n)$ in the worst case
 - How?
 - It is always $\Omega(\log n)$
 - #nodes increases from one level to the next by factor ≤ 2
- Looking for a data structure that supports these operations in time $O(\log n)$ in the worst case
 - Referred to as a *balanced tree* (עץ מאוזן)
- There exist balanced tree DSs based on binary search trees
 - AVL trees, red-black trees
- In this lecture: balanced tree DS that follows a different approach
- Note: the height of a binary search tree is $O(\log n)$ on average
 - What does “on average” mean for a DS?
 - Usually efficient in practice

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
- Dynamic updates
- Discussion

3 Binary heaps

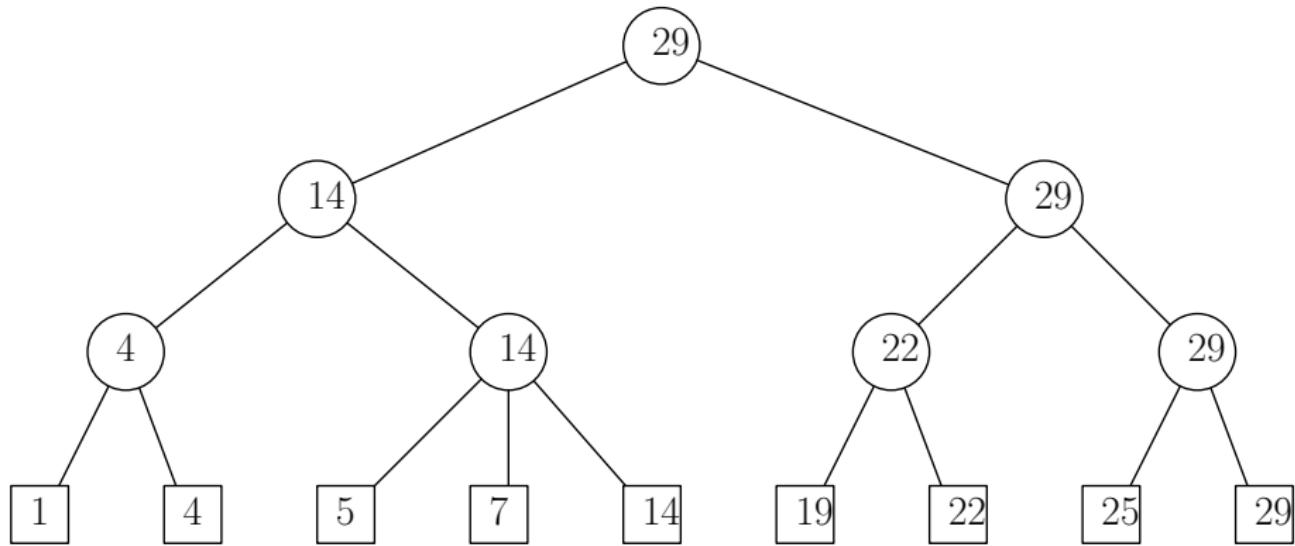
- Binary heap operations
- Discussion

A different kind of rooted tree

2-3 tree (2-3 יער)

- Rooted tree T satisfying:
 - Each internal node has **degree 2 or 3**
 - All leaves are at the **same level**
- Objects are stored **only** at the leaves
 - Internal nodes are there for organization purposes
- Data structure's attribute:
 - **root** = pointer to root of T
- Node attributes:
 - **left** = pointer to left child in T
 - **middle** = pointer to middle child in T
 - **right** = pointer to right child in T
 - **p** = pointer to parent in T
 - **key** = maximum key in its subtree
 - Leaves have additional satellite attributes (storing the actual data)
- When degree is 2, **left** and **middle** children are in use
- The **2-3 tree property**:
keys in left subtree < keys in middle subtree < keys in right subtree

Example



Structural properties

- Consider a 2-3 tree T with n leaves
- Observation 1: in the leaf level, keys are sorted from left to right
 - By induction on height
- Observation 2: #nodes in level $i + 1 \geq 2 \times$ #nodes in level i
 - Each internal node has ≥ 2 children
- Corollary 1: $\text{height}(T) = O(\log n)$
 - Will see: run-time of each supported operation is $O(\text{height}(T))$
- Corollary 2: #internal nodes $<$ #leaves
 - By induction on height
 - Don't waste space (asymptotically)

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations

- Dynamic updates

- Discussion

3 Binary heaps

- Binary heap operations

- Discussion

Initialization

Create an (empty) 2-3 tree

2_3_Init(T)

- 1: new internal node x
- 2: new leaves ℓ, m
- 3: $\ell.key = -\infty$
- 4: $m.key = +\infty$
- 5: $\ell.p = m.p = x$
- 6: $x.key = +\infty$
- 7: $x.left = \ell$
- 8: $x.middle = m$
- 9: $T.root = x$

- ▷ DS attributes initialized to NIL
- ▷ DS attributes initialized to NIL

- Run-time: $O(1)$
- ℓ and m are called *sentinel nodes*

Searching

Search in the 2-3 tree rooted at x for a node whose key is k

`2_3_Search(x, k)`

```
1: if  $x$  is a leaf then
2:   if  $x.key == k$  then
3:     return  $x$ 
4:   else return  $Nil$ 
5: if  $k \leq x.left.key$  then
6:   return 2_3_Search( $x.left, k$ )
7: else if  $k \leq x.middle.key$  then
8:   return 2_3_Search( $x.middle, k$ )
9: else return 2_3_Search( $x.right, k$ )
```

- $+\infty$ sentinel node ensures existence of right child in line 9

Finding the minimum/maximum element

Find the leaf with the smallest key in the 2-3 tree T

2.3 Minimum(T)

```
1:  $x = T.root$ 
2: while  $x$  is not a leaf do
3:    $x = x.left$ 
4:    $x = x.p.middle$ 
5: if  $x.key \neq +\infty$  then
6:   return  $x$ 
7: else error:  $T$  is empty
```

sentinel node, 1, 2, 3, ..., ∞
first less than or equal to p

- How do we find the maximum?

למציאת המינימום ב-2-3 עץ

Finding a successor/predecessor

Find the leaf y with the smallest key among those with $y.key > x.key$

2.3_Successor(x)

```
1:  $z = x.p$ 
2: while  $x == z.right$  or ( $z.right == NIL$  and  $x == z.middle$ ) do
3:    $x = z$            X < y < i נא ש.מ.א. ו.ל.כ. ב.ג.נ.ר.פ.י.ג.ג. X י.ת.ב.
4:    $z = z.p$ 
5: if  $x == z.left$  then
6:    $y = z.middle$ 
7: else  $y = z.right$ 
8: while  $y$  is not a leaf do
9:    $y = y.left$ 
10: if  $y.key < +\infty$  then
11:   return  $y$ 
12: else return  $NIL$ 
```

- How do we find a predecessor?

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
 - Dynamic updates
- Discussion

3 Binary heaps

- Binary heap operations
- Discussion

Procedure Update_Key

Update the key of x to the maximum key in its subtree;
(only) $x.middle$ and $x.right$ may be NIL

Update_Key(x)

- 1: $x.key = x.left.key$
- 2: **if** $x.middle \neq NIL$ **then**
- 3: $x.key = x.middle.key$
- 4: **if** $x.right \neq NIL$ **then**
- 5: $x.key = x.right.key$

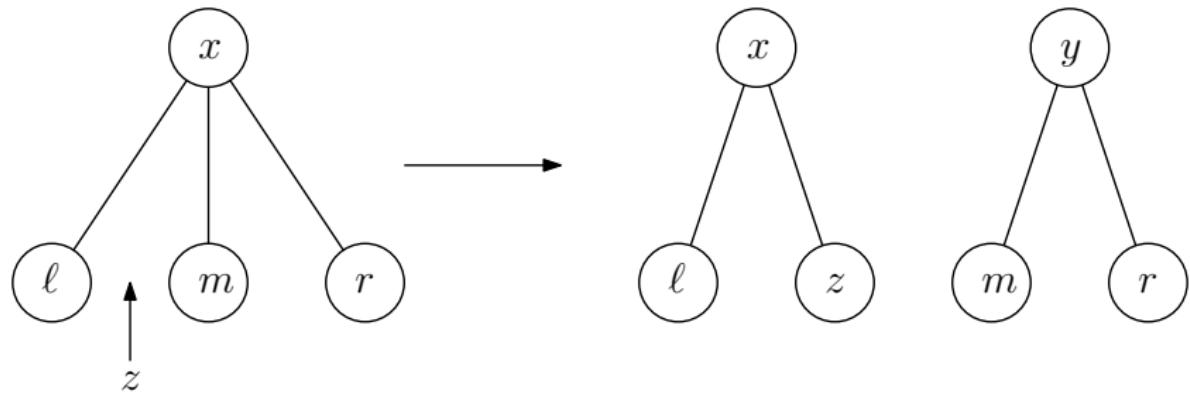
Procedure Set_Children

Set ℓ , m , and r to be the left, middle, and right children, respectively, of x ;
(only) m and r may be NIL

`Set_Children(x, ℓ, m, r)`

- 1: $\langle x.\text{left}, x.\text{middle}, x.\text{right} \rangle = \langle \ell, m, r \rangle$
- 2: $\ell.p = x$
- 3: **if** $m \neq \text{NIL}$ **then**
- 4: $m.p = x$
- 5: **if** $r \neq \text{NIL}$ **then**
- 6: $r.p = x$
- 7: `Update_Key(x)`

New leaf insertion



New leaf insertion — procedure Insert_And_Split

Insert node z as a child of node x ;
split x if necessary and return the new node

Insert_And_Split(x, z)

```
1:  $\langle \ell, m, r \rangle = \langle x.\text{left}, x.\text{middle}, x.\text{right} \rangle$ 
2: if  $r == \text{NIL}$  then
3:   if  $z.\text{key} < \ell.\text{key}$  then
4:     Set_Children( $x, z, \ell, m$ )
5:   else if  $z.\text{key} < m.\text{key}$  then
6:     Set_Children( $x, \ell, z, m$ )
7:   else Set_Children( $x, \ell, m, z$ )
8:   return  $\text{NIL}$ 
9: new internal node  $y$ 
10: ...
```

New leaf insertion — procedure Insert_And_Split

Insert_And_Split(x, z) — cont.

```
10: if  $z.key < \ell.key$  then
11:   Set_Children( $x, z, \ell, NIL$ )
12:   Set_Children( $y, m, r, NIL$ )
13: else if  $z.key < m.key$  then
14:   Set_Children( $x, \ell, z, NIL$ )
15:   Set_Children( $y, m, r, NIL$ )
16: else if  $z.key < r.key$  then
17:   Set_Children( $x, \ell, m, NIL$ )
18:   Set_Children( $y, z, r, NIL$ )
19: else Set_Children( $x, \ell, m, NIL$ )
20:   Set_Children( $y, r, z, NIL$ )
21: return  $y$ 
```

Inserting a new leaf

Insert the new leaf z (DS attributes initialized to NIL) into T

`2_3_Insert(T, z)`

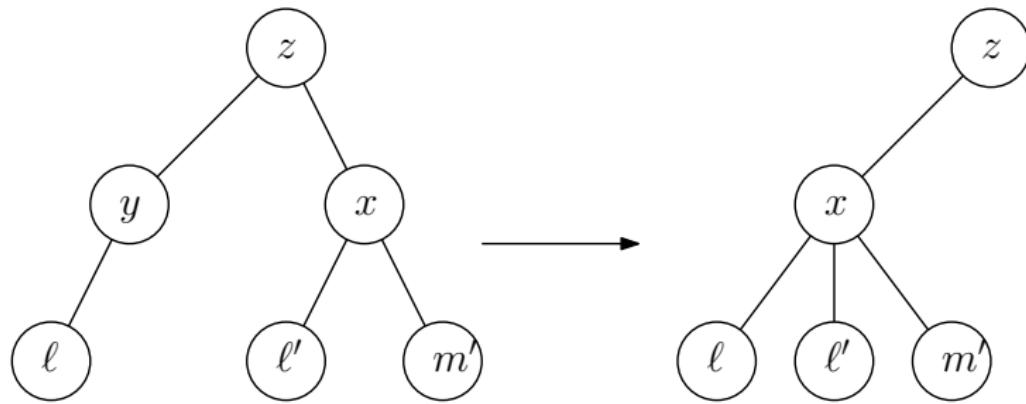
```
1:  $y = T.\text{root}$ 
2: while  $y$  is not a leaf do
3:   if  $z.\text{key} < y.\text{left}.\text{key}$  then  $y = y.\text{left}$ 
4:   else if  $z.\text{key} < y.\text{middle}.\text{key}$  then  $y = y.\text{middle}$ 
5:   else  $y = y.\text{right}$ 
6:  $x = y.p$ 
7:  $z = \text{Insert\_And\_Split}(x, z)$ 
8: while  $x \neq T.\text{root}$  do
9:    $x = x.p$ 
10:  if  $z \neq \text{NIL}$  then
11:     $z = \text{Insert\_And\_Split}(x, z)$ 
12:  else Update_Key( $x$ )
13:  ...
```

Inserting a new leaf

2_3_Insert(T, z) — cont.

```
13: if  $z \neq NIL$  then
14:     new internal node  $w$ 
15:     Set_Children( $w, x, z, NIL$ )
16:      $T.root = w$ 
```

Leaf deletion



Leaf deletion — procedure Borrow_Or_Merge

Borrow a child from a sibling x of y or merge x and y ;
return a pointer to the parent of y (and x)

Borrow_Or_Merge(y)

```
1:  $z = y.p$ 
2: if  $y == z.left$  then
3:    $x = z.middle$ 
4:   if  $x.right \neq NIL$  then
5:     Set_Children( $y, y.left, x.left, NIL$ )
6:     Set_Children( $x, x.middle, x.right, NIL$ )
7:   else Set_Children( $x, y.left, x.left, x.middle$ )
8:   delete  $y$ 
9:   Set_Children( $z, x, z.right, NIL$ )
10:  return  $z$ 
11:  ...
```

Leaf deletion — procedure Borrow_Or_Merge

Borrow_Or_Merge(y) — cont.

```
11: if  $y == z.middle$  then
12:    $x = z.left$ 
13:   if  $x.right \neq NIL$  then
14:     Set_Children( $y, x.right, y.left, NIL$ )
15:     Set_Children( $x, x.left, x.middle, NIL$ )
16:   else Set_Children( $x, x.left, x.middle, y.left$ )
17:   delete  $y$ 
18:   Set_Children( $z, x, z.right, NIL$ )
19: return  $z$ 
20: ...
```

Leaf deletion — procedure Borrow_Or_Merge

Borrow_Or_Merge(y) — cont.

```
20:  $x = z.middle$ 
21: if  $x.right \neq NIL$  then
22:   Set_Children( $y, x.right, y.left, NIL$ )
23:   Set_Children( $x, x.left, x.middle, NIL$ )
24: else Set_Children( $x, x.left, x.middle, y.left$ )
25:   delete  $y$ 
26:   Set_Children( $z, z.left, x, NIL$ )
27: return  $z$ 
```

Deleting a leaf

Delete leaf x from T

2.3 Delete(T, x)

- 1: $y = x.p$
 - 2: **if** $x == y.left$ **then**
 - 3: Set_Children($y, y.middle, y.right, NIL$)
 - 4: **else if** $x == y.middle$ **then**
 - 5: Set_Children($y, y.left, y.right, NIL$)
 - 6: **else** Set_Children($y, y.left, y.middle, NIL$)
 - 7: delete x
 - 8: ...
- ▷ $\deg(y)$ may be < 2

Deleting a leaf

2.3 Delete(T, x) — cont.

```
8: while  $y \neq NIL$  do
9:   if  $y.middle == NIL$  then
10:    if  $y \neq T.root$  then
11:       $y = Borrow\_Or\_Merge(y)$ 
12:    else  $T.root = y.left$ 
13:       $y.left.p = NIL$ 
14:      delete  $y$ 
15:      return
16:    else Update_Key( $y$ )
17:     $y = y.p$ 
```

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
- Dynamic updates

• Discussion

3 Binary heaps

- Binary heap operations
- Discussion

More on 2-3 trees

- All operations run in time $O(\text{height}(T)) = O(\log n)$
 - A balanced tree DS
- Sometimes the actual keys are stored only at the leaves
 - Internal nodes store **pointers** to keys in the corresponding leaves
 - Saving storage if keys themselves are large
- 2-3 trees are a special case of **B^+ trees**
 - Parameterized by a **degree parameter d**
 - The degree of each internal node x satisfies $\lceil d/2 \rceil \leq \deg(x) \leq d$
 - The degree of the root r satisfies $2 \leq \deg(r) \leq d$
 - Height proportional to $\log_d n$
- Often: x stores maximum key in **each child subtree**
 - $\deg(x)$ keys in total
- Very efficient for **external memories**
 - Access is very expensive in comparison to internal memory
 - Read/write a whole **page** rather than a single address
 - Adjust the parameter d so that size of a node \approx size of page
 - Save on #page read/write operations

1 Binary search trees

- Binary search tree operations

2 2-3 trees

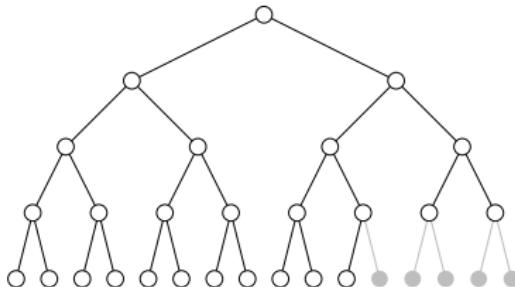
- 2-3 tree operations
- Dynamic updates
- Discussion

3 Binary heaps

- Binary heap operations
- Discussion

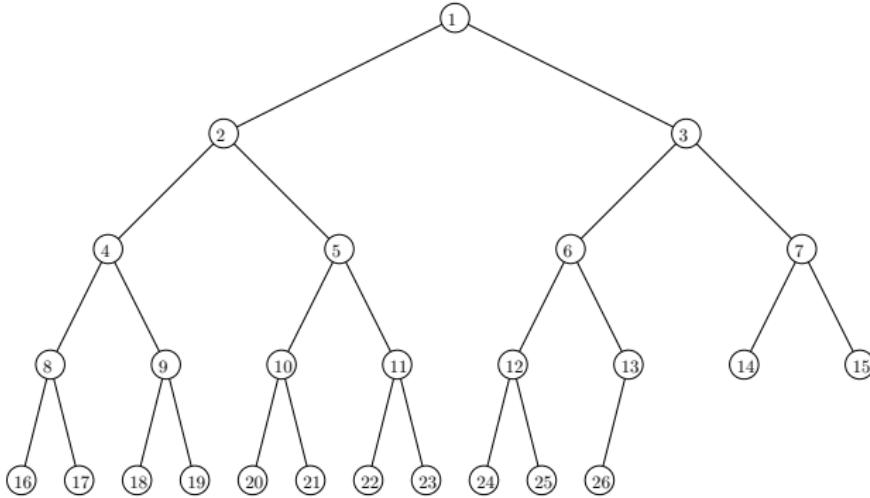
Binary trees revisited

- A *complete* (שלם) binary tree:
 - the degree of each internal node is exactly 2; and
 - all leaves are at the same depth
- Height h complete binary tree has 2^h leaves and $2^h - 1$ internal nodes
- A *nearly complete* (כמעטם שלם) binary tree:
 - obtained from a height h complete binary tree by deleting the k rightmost leaves for some $0 \leq k < 2^h$
- Properties of nearly complete binary trees:
 - A generalization of complete binary trees
 - All leaves are in the last two levels
 - $\#\text{internal nodes} \leq \#\text{leaves} \leq \#\text{internal nodes} + 1$



Array representation

- Representing an n -node nearly complete binary tree in array $A[1 \dots n]$
 - The root is identified with $A[1]$
 - Left child of internal node $A[i]$ is identified with $A[2i]$
 - Right child of internal node $A[i]$ is identified with $A[2i + 1]$
 - Parent of non-root node $A[i]$ is identified with $A[\lfloor i/2 \rfloor]$
 - Define $\text{Left}(i) = 2i$, $\text{Right}(i) = 2i + 1$, $\text{Parent}(i) = \lfloor i/2 \rfloor$



A nearly complete binary tree data structure

ערמה בינארית ⟨Binary heap⟩

(sometimes simply heap)

- Objects associated with the nodes of a nearly complete binary tree T
- T represented in (prefix of) array $A[1 \dots A.length]$
 - $A.length$ is an upper bound on #objects in the dynamic set
 - Known in advance
 - Exists in many applications
 - n objects stored in $A[1 \dots n]$
- Data structure's attribute (on top of the standard array's attributes):
 - $heap-size = \#$ currently stored objects ($n \leq A.length$)
- The binary heap property:
 $A[\text{Parent}(i)].key < A[i].key$ for every non-root $A[i]$ each almost complete binary tree that have this is a binary heap
- Sometimes referred to as a priority queue or **minimum heap** **ערמת מינימום** **ערמת מקסימום**
 - Switching the relevant operations/inequalities yields a maximum heap **ערמת מקסימום**

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
- Dynamic updates
- Discussion

3 Binary heaps

- Binary heap operations
- Discussion

Building a heap — Procedure Heapify

Precondition: subtrees rooted at $A[\text{Left}(i)]$ and $A[\text{Right}(i)]$ are heaps;
modify A so that the subtree rooted at $A[i]$ is a heap

Heapify(A, i)

```
1:  $\ell = \text{Left}(i)$ 
2: if  $\ell \leq A.\text{heap-size}$  and  $A[\ell].\text{key} < A[i].\text{key}$  then
3:    $\text{smallest} = \ell$ 
4: else  $\text{smallest} = i$ 
5:  $r = \text{Right}(i)$ 
6: if  $r \leq A.\text{heap-size}$  and  $A[r].\text{key} < A[\text{smallest}].\text{key}$  then
7:    $\text{smallest} = r$ 
8: if  $\text{smallest} \neq i$  then
9:   swap  $A[i]$  and  $A[\text{smallest}]$ 
10:  Heapify( $A, \text{smallest}$ )
```

- $A[i]$ **seeps down** the tree
- Run-time: $O(\text{height}(T_i)) = O(\log n)$
 - T_i = subtree rooted at $A[i]$

Building a heap from an arbitrary array

Build a heap from the objects in $A[1 \dots A.length]$

`Build_Heap(A)`

- 1: $A.heap-size = A.length$
- 2: **for** $i = A.length, \dots, 2, 1$ **do**
- 3: `Heapify(A, i)`

- Naive run-time analysis:
 - n iterations
 - $O(\log n)$ time per iteration
 - $\Rightarrow O(n \log n)$
- Can do better

Building a heap — run-time analysis

- Recall: run-time of `Heapify(A, i)` is $O(\text{height}(T_i))$
- Height of a nearly complete binary tree T with n nodes is $\lfloor \lg n \rfloor$
- #nodes of height j in $T \leq 2^{\lfloor \lg n \rfloor - j} = O(n/2^j)$
- \Rightarrow run-time of `Build_Heap` is

$$O\left(\sum_{j=0}^{\lfloor \lg n \rfloor} j \cdot \frac{n}{2^j}\right) \leq O(n) \cdot \sum_{j=0}^{\infty} \frac{j}{2^j} = O(n)$$

Extracting the minimum

Delete the minimum object from heap A and return it

Heap_Extract_Min(A)

- 1: **if** $A.\text{heap-size} < 1$ **then**
- 2: error “the heap is empty”
- 3: $\min = A[1]$
- 4: $A[1] = A[A.\text{heap-size}]$
- 5: $A.\text{heap-size} = A.\text{heap-size} - 1$
- 6: Heapify($A, 1$)
- 7: **return** \min

- Run-time: $O(\log n)$
- What's the run-time of returning (a pointer to) the minimum object without deleting it? $O(1)$

Decreasing a key

Precondition: $k < A[i].key$;

change the key of $A[i]$ to k

`HeapDecreaseKey(A, i, k)`

- 1: **if** $k > A[i].key$ **then**
- 2: error “new key is larger than current key”
- 3: $A[i].key = k$
- 4: **while** $i > 1$ and $A[i].key < A[\text{Parent}(i)].key$ **do**
- 5: swap $A[i]$ and $A[\text{Parent}(i)]$
- 6: $i = \text{Parent}(i)$

- The updated $A[i]$ **seeps up** the tree
- Run-time: $O(\log n)$

Inserting a new node

Insert the new heap node x into heap A

`Heap_Insert(A, x)`

- 1: $s = A.\text{heap-size} + 1$
- 2: $A[s] = x$ ▷ copy all satellite attributes
- 3: $A[s].\text{key} = \infty$
- 4: $A.\text{heap-size} = s$ ▷ a valid heap of size s
- 5: `Heap_Decrease_Key($A, s, x.\text{key}$)`

- Run-time: $O(\log n)$

1 Binary search trees

- Binary search tree operations

2 2-3 trees

- 2-3 tree operations
- Dynamic updates
- Discussion

3 Binary heaps

- Binary heap operations
- Discussion

Comparison to balanced trees

- **Balanced trees** can be adapted to support:
 - Returning a pointer to the minimum object in time $O(1)$
 - Extracting the minimum object in time $O(\log n)$
 - Decreasing the key of a given node in time $O(\log n)$
 - How?
- **Advantages of heaps over balanced trees:**
 - Building a new heap with n objects is (asymptotically) faster
 - Speed-up (non-asymptotic) due to smaller constants
 - Speed-up (non-asymptotic) due to array representation
 - No pointers, whole DS stored in the same memory page