

Shortest Paths in Weighted Graphs

Data Structures and Algorithms (094224)

Yuval Emek

Winter 2022/23

1 Weighted graphs

2 Foundations of single source shortest paths algorithms

- Relaxation
- Properties of SSSP

3 Bellman-Ford's algorithm

- Correctness

4 Dijkstra's algorithm

- Correctness
- Run-time

5 All pairs shortest paths: Floyd-Warshall's algorithm

- Performance

Distances revisited

- Digraph $G = (V, E)$
 - Some of the material in this lecture applies to undirected graphs as well
- Edge **weight** $\langle \text{משקל} \rangle$ function $w : E \rightarrow \mathbb{R}$
 - Associates a real weight $w(e)$ with each edge $e \in E$
 - G is called a **weighted** $\langle \text{ ממוקל} \rangle$ digraph
- Generalize to edge subsets $F \subseteq E$: $w(F) = \sum_{e \in F} w(e)$
- Weight of **path** $P = \langle v_0, v_1, \dots, v_k \rangle$ is $w(P) = \sum_{i=1}^k w(v_{i-1}, v_i)$
 - A.k.a. (weighted) **length** of P
- Redefine the **distance** from $u \in V$ to $v \in V$ w.r.t. w :
$$\delta(u, v) = \begin{cases} \min \left\{ w(P) : u \xrightarrow{P} v \right\}, & u \rightsquigarrow v \\ \infty, & u \not\rightsquigarrow v \end{cases}$$
 - $u \rightsquigarrow v$ denotes v reachable from u
 - $u \xrightarrow{P} v$ denotes path P from u to v
- P is a **shortest path** from u to v if $w(P) = \delta(u, v)$
- The **unweighted** $\langle \text{ ממוקל לא} \rangle$ case: $w(e) = 1$ for every $e \in E$

Negative weights

- Are the distances well defined in the presence of **negative** weights?
- What if G admits a **negative weight cycle** C ?
 - $\delta(u, v)$ is **unbounded** (from below) for every $u \rightsquigarrow C \rightsquigarrow v$
 - Can make the path arbitrarily short by including more traversals of C
 - Define $\delta(u, v) = -\infty$
- Some shortest paths algorithms assume that the edge weights are **non-negative**
- Other shortest paths algorithms allow negative weight edges but **forbid** negative weight cycles
- If G doesn't admit negative weight cycles, then we can restrict our attention to **simple** shortest paths
 - Does it mean that every shortest path is necessarily simple?
yes, but there can be (kayyam) not simple path if there is a cycle with 0 weight (we can still find shortest simple path).

Subpaths of shortest paths are shortest

this lemma is true to (di/un)graphs

Lemma (subpaths of shortest paths)

Consider a shortest path $P = \langle v_0, v_1, \dots, v_k \rangle$. Then, $P_{i,j} = \langle v_i, \dots, v_j \rangle$ is a shortest path for every $0 \leq i \leq j \leq k$.

- Assume by contradiction: $w(Q) < w(P_{i,j})$ for some (v_i, v_j) -path Q
- Let $P' = v_0 \xrightarrow{P_{0,i}} v_i \xrightarrow{Q} v_j \xrightarrow{P_{j,k}} v_k$
- $w(P') = w(P_{0,i}) + w(Q) + w(P_{j,k}) = w(P) - w(P_{i,j}) + w(Q) < w(P)$
- P is not a shortest (v_0, v_k) -path ($\rightarrow\leftarrow$) ■

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

Representing shortest paths

- Focus: single source shortest paths (**SSSP**) algorithm
 - Designated **source** vertex $s \in V$
 - **Goal:** compute the distances and shortest paths from s to all vertices reachable from s
 - BFS does so for unweighted (di)graphs
- A **shortest paths tree**
 - Tree rooted at s
 - Contains all vertices reachable from s
 - Unique simple (s, v) -path in the tree is a shortest (s, v) -path in G

The structure of the algorithms

- The SSSP algorithms we will see have many common features
- Additional attributes for each vertex $v \in V$:
 - $v.d$ = the distance from s to v in G
 - $v.\pi$ = the parent of v in the shortest paths tree
- The (directed) predecessor subgraph $G_\pi = (V_\pi, E_\pi)$
 - $V_\pi = \{v \in V \mid v.\pi \neq NIL\} \cup \{s\}$
 - $E_\pi = \{(v.\pi, v) \in E \mid v \in V_\pi - \{s\}\}$
 - At termination: (its undirected version is) a shortest paths tree for s
- Algorithm's structure:
 - ① Initialize the d and π fields by calling procedure
`Initialize_Single_Source`
 - ② Update the d and π fields through calls to procedure `Relax`
 - d and π are **not modified otherwise**

The initialization procedure

Initialize the d and π fields in G

`Initialize_Single_Source(G, s)`

- 1: **for all** $v \in G.V$ **do**
- 2: $v.d = \infty$
- 3: $v.\pi = NIL$
- 4: $s.d = 0$

- Run-time: $O(n)$

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

The triangle inequality

Lemma (triangle inequality)

For every edge $(u, v) \in E$, we have $\delta(s, v) \leq \delta(s, u) + w(u, v)$.

- Case $\delta(s, v) = \infty$: $s \not\rightsquigarrow v$, thus $s \not\rightsquigarrow u$ and $\delta(s, u) = \infty$
- Case $\delta(s, u) = -\infty$: there exists a negative weight cycle C such that $s \rightsquigarrow C \rightsquigarrow u$, thus $s \rightsquigarrow C \rightsquigarrow v$ and $\delta(s, v) = -\infty$
- Assume that $-\infty < \delta(s, u), \delta(s, v) < \infty$
- Let Q be a shortest (s, u) -path
- Let $P = s \stackrel{Q}{\rightsquigarrow} u \rightarrow v$
- $w(P) = w(Q) + w(u, v) = \delta(s, u) + w(u, v)$
- P is a candidate for a shortest (s, v) -path, thus $\delta(s, v) \leq w(P)$ ■

Procedure Relax

Try to improve the shortest path to v through the edge (u, v)

Relax(u, v, w)

- 1: **if** $v.d > u.d + w(u, v)$ **then**
- 2: $v.d = u.d + w(u, v)$ triangle inequality
- 3: $v.\pi = u$

- Run-time: $O(1)$
- Improvements based on “realizing the triangle inequality”
- The basic operation of our algorithms
 - The only means by which d and π are modified after initialization

1 Weighted graphs

2 Foundations of single source shortest paths algorithms

- Relaxation
- Properties of SSSP

3 Bellman-Ford's algorithm

- Correctness

4 Dijkstra's algorithm

- Correctness
- Run-time

5 All pairs shortest paths: Floyd-Warshall's algorithm

- Performance

The monotonicity property

Lemma (monotonicity)

The value of $v.d$ is non-increasing over time.

- Set to $+\infty$ during initialization
- Can only decrease during calls to procedure Relax ■

The upper bound property

Lemma (upper bound)

$v.d \geq \delta(s, v)$ at all times.

- Proof by induction on the sequence of calls to procedure Relax
- Base: immediately after initialization, $v.d = \infty \geq \delta(s, v)$ for all $v \in V - \{s\}$ and $s.d = 0 \geq \delta(s, s)$
- Step: consider some call to $\text{Relax}(u, v, w)$
- The only d value that may change during the call is $v.d$
- If it changes, then after the call

$$\begin{aligned} v.d &= u.d + w(u, v) \\ &\geq \delta(s, u) + w(u, v) && \text{ind. hyp.} \\ &\geq \delta(s, v) && \text{triangle inequality } \blacksquare \end{aligned}$$

The no path property

Corollary (no path)

If v is not reachable from s , then $v.d = \infty$ at all times.

- By the upper bound property, $v.d \geq \delta(s, v) = \infty$ at all times ■

The convergence property

Lemma (convergence)

If $s \rightsquigarrow u \rightarrow v$ is a shortest (s, v) -path and $u.d = \delta(s, u)$ prior to some call to $\text{Relax}(u, v, w)$, then $v.d = \delta(s, v)$ at all times after the call.

- Right after the call

$$\begin{aligned}v.d &\leq u.d + w(u, v) \\&= \delta(s, u) + w(u, v) && \text{assumption on } u.d \\&= \delta(s, v) && s \rightsquigarrow u \rightarrow v \text{ is shortest}\end{aligned}$$

- By the upper bound property, inequality cannot be strict
- Subsequently, the value of $v.d$ doesn't change by the monotonicity and upper bound properties ■

The path relaxation property

Lemma (path relaxation)

Consider some path $P = \langle s = v_0, v_1, \dots, v_k \rangle$. If the sequence of calls to Relax includes as a subsequence^a the calls $\text{Relax}(v_0, v_1, w), \text{Relax}(v_1, v_2, w), \dots, \text{Relax}(v_{k-1}, v_k, w)$, then $v_k.d \leq w(P)$ at all times after this subsequence of calls.

^aNot necessarily contiguous.

- Prove by induction on $i = 0, 1, \dots, k$ that after edge (v_{i-1}, v_i) is relaxed, $v_i.d \leq \sum_{j=1}^i w(v_{j-1}, v_j)$
- **Base:** after initialization, $v_0.d = s.d = 0$ and this value doesn't increase by the monotonicity property
- **Step:** consider the call to $\text{Relax}(v_i, v_{i+1}, w)$
 - By the ind. hyp., $v_i.d \leq \sum_{j=1}^i w(v_{j-1}, v_j)$ before the call
 - Relax ensures that $v_{i+1}.d \leq v_i.d + w(v_i, v_{i+1})$ right after the call
 - Doesn't increase by the monotonicity property ■

The shortest path relaxation property

Corollary (shortest path relaxation)

Consider some shortest path $P = \langle s = v_0, v_1, \dots, v_k \rangle$. If the sequence of calls to Relax includes as a subsequence^a the calls

$\text{Relax}(v_0, v_1, w), \text{Relax}(v_1, v_2, w), \dots, \text{Relax}(v_{k-1}, v_k, w)$, then
 $v_k.d = \delta(s, v_k)$ at all times after this subsequence of calls.

^aNot necessarily contiguous.

- By the path relaxation property, $v_k.d \leq w(P) = \delta(s, v_k)$ at all times after the subsequence
- The assertion follows by the upper bound property ■

The predecessor subgraph

Lemma

If G is free of negative weight cycles reachable from s , then the undirected version of G_π forms a tree at all times.

Prove:

- ① G_π is cycle free
- ② The undirected version of G_π is cycle free
- ③ If $v \in V_\pi$, then v is reachable from s in G_π

G_π is cycle free

- Assume by contradiction: G_π admits cycle $C = \langle v_0, v_1, \dots, v_k = v_0 \rangle$
- W.l.o.g.: (v_{k-1}, v_k) is the last edge of the cycle to join E_π
 - Happened by relaxing the edge at some time t^* the time C was created in G_pi
- Claim 1: $s \rightsquigarrow C$
 - Since $v_i.\pi \neq NIL$, we know that $v_i.d < \infty$, hence $s \rightsquigarrow v_i$ by the no path property
- Claim 2: right before time t^* , $v_i.d \geq v_{i-1}.d + w(v_{i-1}, v_i)$ for every $1 \leq i \leq k - 1$
 - $v_i.\pi$ was set to v_{i-1} during a call to $\text{Relax}(v_{i-1}, v_i, w)$ before time t^*
 - Right after this call, we had $v_i.d = v_{i-1}.d + w(v_{i-1}, v_i)$
 - $v_i.d$ didn't change since then because $v_i.\pi$ didn't
 - By the monotonicity property, since then, $v_{i-1}.d$ could only decrease
- Claim 3: right before time t^* , $v_k.d > v_{k-1}.d + w(v_{k-1}, v_k)$
 - The if condition (line 1) was satisfied during the call at time t^*

G_π is cycle free — cont.

- Combining Claims 2 and 3: right before time t^* ,

$$\sum_{i=1}^k v_i.d > \sum_{i=1}^k (v_{i-1}.d + w(v_{i-1}, v_i)) = \sum_{i=1}^k v_{i-1}.d + \sum_{i=1}^k w(v_{i-1}, v_i)$$

- Since $\sum_{i=1}^k v_i.d = \sum_{i=1}^k v_{i-1}.d$, we get

$$w(C) = \sum_{i=1}^k w(v_{i-1}, v_i) < 0$$

- ⇒ A negative weight cycle reachable from s (by Claim 1) ($\rightarrow\leftarrow$)

The undirected version of G_π is cycle free

- Assume by contradiction:
 - undirected version of G_π admits a simple cycle C
 - $k = \#\text{vertices in } C (= \#\text{edges in } C)$
- Consider the digraph obtained from C by directing the edges according to G_π
- For every vertex v , $\deg_{\text{in}}(v) + \deg_{\text{out}}(v) = 2$
- Sum of in-degrees = sum of out-degrees = k
 - \Rightarrow Average in-degree = average out-degree = 1
- C is not a (directed) cycle in G_π , hence not all vertices v satisfy $\deg_{\text{in}}(v) = \deg_{\text{out}}(v) = 1$
 - \Rightarrow There must exist a vertex v with $\deg_{\text{in}}(v) > 1$
- In contradiction to the construction of G_π ($\rightarrow \leftarrow$)

If $v \in V_\pi$, then v is reachable from s in G_π

- Starting with $v = v_0$, we tour the graph by moving from v_i to $v_{i+1} = v_i.\pi$
- This tour must stop at some vertex v_k with $v_k.\pi = NIL$
 - Otherwise, we reached a (directed) cycle in G_π — proved already that this can't be
- Assume by contradiction: $v_k \neq s$
- $v_{k-1}.\pi$ was set to v_k by calling $\text{Relax}(v_k, v_{k-1}, w)$
- At the time of this call $v_k.d$ must have been $< \infty$
- \Rightarrow At that time, $v_k.\pi$ must have been $\neq NIL$ and it remained $\neq NIL$ ever since ($\rightarrow\leftarrow$)
- Therefore $v_k = s$ and v is reachable from s in G_π ■

The shortest paths property

Lemma (shortest paths)

Suppose that G is free of negative weight cycles reachable from s . At all times, if $v.d = \delta(s, v) < \infty$, then the unique (s, v) -path in G_π is a shortest (s, v) -path in G .

- Let t^v be the time at which $v.d$ is set to $v.d = \delta(s, v)$
- By the monotonicity and upper bound properties, $v.d$ is not updated after time t^v , hence $v.\pi$ is not updated after time t^v
- Claim:** if $v.\pi$ is set to u at time t^v , then $t^u < t^v$ and $\delta(s, v) = \delta(s, u) + w(u, v)$
 - Relax(u, v, w) is called at time t^v and at that time $v.d$ is set to

$$\begin{aligned}\delta(s, v) &= v.d = u.d + w(u, v) \\ &\geq \delta(s, u) + w(u, v) && \text{upper bound property} \\ &\geq \delta(s, v) && \text{triangle inequality}\end{aligned}$$

- the inequalities are not strict, establishing the claim

The shortest paths property — cont.

- Let $P_v(t)$ be the unique (s, v) -path in G_π at time t
 - Assuming that $v \in V_\pi$ at time t
- Prove by induction on the times t^v : $\text{len}(P_v(t^v)) = \delta(s, v)$ and $P_v(t) = P_v(t^v)$ for all $t \geq t^v$
- **Base:** the empty path is the unique (s, s) -path in G_π at all times and its length is indeed 0
- **Step:** consider some v whose $v.\pi$ field is set to $v.\pi = u$ at time t^v
- By the claim, $t^u < t^v$ and $\delta(s, v) = \delta(s, u) + w(u, v)$
- By the ind. hyp., $\text{len}(P_u(t^u)) = \delta(s, u)$ and $P_u(t) = P_u(t^u)$ for all $t \geq t^u$

The shortest paths property — cont.

- Hence,

$$\begin{aligned}\text{len}(P_v(t^v)) &= \text{len}(P_u(t^v)) + w(u, v) \\ &= \text{len}(P_u(t^u)) + w(u, v) && t^v > t^u \\ &= \delta(s, u) + w(u, v) \\ &= \delta(s, v)\end{aligned}$$

- $v.\pi$ is not updated after time t^v , therefore $P_v(t) = P_v(t^v)$ for all $t \geq t^v$ ■

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

The Bellman-Ford algorithm


- Input:
 - Digraph $G = (V, E)$
 - Weight function $w : E \rightarrow \mathbb{R}$
 - Source vertex $s \in V$
- Outputs an **error message** if G admits a negative weight cycle reachable from s
- Computes (if no error):
 - $\delta(s, v)$ for each vertex $v \in V$
 - Shortest paths tree rooted at s

Pseudocode

`Bellman_Ford(G, w, s)`

```
1: Initialize_Single_Source( $G, s$ )
2: for  $i = 1, \dots, |G.V| - 1$  do
3:   for all  $(u, v) \in G.E$  do
4:     Relax( $u, v, w$ )
5:   for all  $(u, v) \in G.E$  do
6:     if  $v.d > u.d + w(u, v)$  then
7:       error: “negative weight cycle”
```

Example



Run-time analysis

- The initialization takes $O(n)$ time
- The for loop of lines 2–4 makes $O(n)$ iterations
- The for loops of lines 3–4 and lines 5–7 make $O(m)$ iterations
- Each call to Relax takes $O(1)$ time
- $O(mn)$ time in total
- **Exercise:** show that this is asymptotically tight

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

Correctness of the distances

Lemma

If G is free of negative weight cycles reachable from s , then upon termination, $v.d = \delta(s, v)$ for every vertex $v \in V$.

- Holds for vertices v not reachable from s by the no path property
- Consider a simple shortest (s, v) -path $\langle s = v_0, v_1, \dots, v_k = v \rangle$
 - Simple path: $k \leq n - 1$
- Edge (v_{i-1}, v_i) is relaxed in iteration i of the for loop of lines 2–4
- The assertion follows by the shortest path relaxation property ■

Correctness of the shortest paths tree

Corollary

If G is free of negative weight cycles reachable from s , then upon termination, (the undirected version of) G_π is a shortest paths tree rooted at s .

- Combine the previous lemma with the shortest paths property ■

Detecting negative weight cycles

Lemma

`Bellman_Ford(G, s, w)` outputs an error message iff G admits a negative weight cycle reachable from s .

- Suppose that G is free of negative weight cycles reachable from s
- **Already proved:** upon termination of the for loop of lines 2–4,
 $v.d = \delta(s, v)$ for every $v \in V$
- Therefore, for every edge $(u, v) \in E$,

$$\begin{aligned} v.d &= \delta(s, v) \\ &\leq \delta(s, u) + w(u, v) && \text{triangle inequality} \\ &= u.d + w(u, v) \end{aligned}$$

- \implies The if condition in line 6 fails for all edges

Detecting negative weight cycles — cont.

- Suppose that G admits a negative weight cycle
 $C = \langle v_0, v_1, \dots, v_k = v_0 \rangle$ reachable from s
- $s \rightsquigarrow v_i$ implies the existence of a simple path $s \xrightarrow{P} v_i$ for every $1 \leq i \leq k$
 - P contains $\leq n - 1$ edges
- By the path relaxation property, upon termination of the for loop of lines 2–4, $v_i.d < \infty$ for every $1 \leq i \leq k$
- Assume by contradiction: if condition in line 6 fails for all edges in C
 - \implies upon termination of the for loop of lines 2–4,
 $v_i.d \leq v_{i-1}.d + w(v_{i-1}, v_i)$ for every $1 \leq i \leq k$
- Summing over all i s: $\sum_{i=1}^k v_i.d \leq \sum_{i=1}^k v_{i-1}.d + \sum_{i=1}^k w(v_{i-1}, v_i)$
- Since $\sum_{i=1}^k v_i.d = \sum_{i=1}^k v_{i-1}.d$, we conclude that
 $\sum_{i=1}^k w(v_{i-1}, v_i) \geq 0$ ($\rightarrow \leftarrow$) ■

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

The Dijkstra algorithm

- **Input:**

- Digraph $G = (V, E)$
- Weight function $w : E \rightarrow \mathbb{R}_{\geq 0}$
 - Note the **non-negative** edge weights
- Source vertex $s \in V$

- **Computes:**

- $\delta(s, v)$ for each vertex $v \in V$
- Shortest paths tree rooted at s

- Works also for **undirected** graphs


Pseudocode

Dijkstra(G, w, s)

```
1: Initialize_Single_Source( $G, s$ )
2:  $Q = G.V$ 
3: while  $Q \neq \emptyset$  do
4:    $u = \text{Extract\_Min}(Q)$                                  $\triangleright$  minimum w.r.t.  $d$ 
5:   for all  $v \in G.\text{Adj}[u]$  do
6:     Relax( $u, v, w$ )
```

- Resemblance to BFS

Example



- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

Correctness of the distances

Lemma

When vertex u is extracted from Q (line 4), it holds that $u.d = \delta(s, u)$.

- Let t^v be the time just before $v \in V$ is extracted from Q (line 4)
- Assume by contradiction: the lemma is false for some vertex $u \in V$
- Let u be the vertex that minimizes t^u among all vertices with $u.d \neq \delta(s, u)$ at time t^u
 - $u.d > \delta(s, u)$ at time t^u by the upper bound property
- $u \neq s$ because s is the first to be extracted from Q with $s.d = 0 = \delta(s, s)$
- $s \leadsto u$ because otherwise $\delta(s, u) = \infty \not\prec u.d$
- Let P be a shortest (s, u) -path

Correctness of the distances — cont.

- At time t^u , P leads from a vertex in $V - Q$ to a vertex in Q
 - Decompose $P = s \rightsquigarrow x \rightarrow y \rightsquigarrow u$ so that $x \notin Q$ and $y \in Q$ at time t^u
- Since $x \notin Q$ at time t^u , it follows that $t^x < t^u$
- By the choice of u , we know that $x.d = \delta(s, x)$ at time t^x
- Edge (x, y) is relaxed soon after time t^x , so $y.d = \delta(s, y)$ at all times afterwards by the convergence property
- At time t^u , we have

$$\begin{array}{ll} y.d = \delta(s, y) & t^u > t^x \\ \leq \delta(s, u) & y \text{ precedes } u \text{ along } P \\ < u.d & \text{assumption on } u \end{array}$$

- $\implies y$ should have been selected in line 4 at time t^u ($\rightarrow \leftarrow$) ■

Correctness of the shortest paths tree

Corollary

Upon termination, (the undirected version of) G_π is a shortest paths tree rooted at s .

- Combine the previous lemma with the shortest paths property ■

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

Run-time analysis

- The initialization takes $O(n)$ time
- Each vertex is extracted from Q (line 4) exactly once
- \Rightarrow Each edge is relaxed (line 6) exactly once
- A naive implementation:
 - Each call to Extract_Min takes $O(n)$ time
 - Run-time: $O(n^2 + m) = O(n^2)$
- Implementing Q using a binary heap:
 - Each call to Extract_Min takes $O(\log n)$ time
 - Each call to Decrease_Key (hidden in line 6) takes $O(\log n)$ time
 - Run-time: $O((n + m) \log n)$
- Implementing Q using a Fibonacci heap:
 - Run-time: $O(n \log n + m)$
 - Beyond the scope of this course

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

Computing all pairs shortest paths

- **Goal:** given a digraph $G = (V, E)$ with edge weight function $w : E \rightarrow \mathbb{R}$, compute the distance (and a shortest path) from u to v for **all** $(u, v) \in V \times V$
- **No negative weight edges:** invoke Dijkstra for every source vertex
 - Run-time: $O(n^2 \log n + nm)$
- **Negative weight edges:** invoke Bellman-Ford for every source vertex
 - Run-time: $O(n^2 m)$
 - Can do better with **dynamic programming**...

The Floyd-Warshall algorithm


- Input:
 - Digraph $G = (V, E)$
 - $V = [n] = \{1, \dots, n\}$
 - Weight function $w : E \rightarrow \mathbb{R}$
 - No negative weight cycles
 - $w(i, j) = \infty$ if $(i, j) \notin E$
- Computes:
 - The **distance** $\delta(i, j)$ for every $i, j \in [n]$
 - The **successor** $s(i, j)$ of i along a shortest (i, j) -path for every $i, j \in [n]$

Intermediate vertices

- v_i is an *intermediate vertex* of path $\langle v_1, \dots, v_k \rangle$ if $1 < i < k$
- Let $P_k(i, j)$ be a shortest (i, j) -path that uses **only** vertices in $[k]$ as intermediate vertices
 - Not necessarily all vertices in $[k]$
- Define $\delta_k(i, j) = w(P_k(i, j))$
- Recursive formula:

$$\delta_k(i, j) = \begin{cases} w(i, j), & k = 0 \\ \min \{\delta_{k-1}(i, j), \delta_{k-1}(i, k) + \delta_{k-1}(k, j)\}, & k > 0 \end{cases}$$

- Construct $P_k(i, j)$ by taking $P_{k-1}(i, j)$ or $P_{k-1}(i, k) \circ P_{k-1}(k, j)$



Pseudocode

Floyd – Warshall(G, w)

```
1:  $n = |G.V|$ 
2: new  $n \times n$  table  $D_0[1 \dots n, 1 \dots n]$ 
3: initialize  $D_0$  so that  $D_0[i, j] = w(i, j)$ 
4: for  $k = 1, \dots, n$  do
5:     new  $n \times n$  table  $D_k[1 \dots n, 1 \dots n]$ 
6:     for  $i = 1, \dots, n$  do
7:         for  $j = 1, \dots, n$  do
8:              $D_k[i, j] = \min \{D_{k-1}[i, j], D_{k-1}[i, k] + D_{k-1}[k, j]\}$ 
9: return  $D_n$ 
```

Retrieving the actual paths

- The aforementioned pseudocode computes the distances (optimal values), but not the shortest paths (optimal solutions)
- Interested in the *successor matrix* $S \in V^{n \times n}$ so that $S[i, j] = s(i, j)$
- Can be constructed by modifying the algorithm to keep track of table updates
 - Often the case in dynamic programming algorithms
- S can also be computed from D_n in time $O(n^3)$ in a **black-box** fashion
 - Exercise

- 1 Weighted graphs
- 2 Foundations of single source shortest paths algorithms
 - Relaxation
 - Properties of SSSP
- 3 Bellman-Ford's algorithm
 - Correctness
- 4 Dijkstra's algorithm
 - Correctness
 - Run-time
- 5 All pairs shortest paths: Floyd-Warshall's algorithm
 - Performance

Run-time analysis

- Initialization takes $O(n^2)$ time
- Line 8 takes $O(1)$ time
- Three nested for loops, each with n iterations
- Total run-time: $O(n^3)$

Space analysis

- The algorithm allocates n tables, each of size $n \times n$
 - $O(n^3)$ space
- Notice: during iteration k of the outer-most for loop, we only need table D_{k-1}
 - Tables D_1, \dots, D_{k-2} can be deleted
- Space decreases to $O(n^2)$
 - Is it asymptotically optimal?