

Efficient Data Structures

Data Structures and Algorithms (094224)

Tutorial 6

Winter 2022/23

1 Binary search trees

2 Augmenting Data Structures

3 Non Distinct Keys

Question 1

A *labeled binary tree* is a binary tree with a unique label in each node

Prove/Disprove:

Let T be a binary tree with n nodes and K a set of n distinct keys. There exists a single mapping of keys from K to nodes in T such that the resulting labeled binary tree maintains the binary search tree property.

The binary search tree property for labeled binary tree:

if y is a node in the left (resp., right) subtree of node x , then
 $label(y) < label(x)$ (resp., $label(y) > label(x)$)

Solution

Definition:

- The i th *order statistic* סטטיסטי הסדר \langle of a set of n elements is the i th smallest element
 - The minimum is the first order statistic ($i = 1$)
 - The maximum is the n th order statistic ($i = n$)
- Example:
 - $S = \{10, 1, 5, 18, 7, 3, 4\}$. The linear order is $(1, 3, 4, 5, 7, 10, 18)$
 - The second order statistic is 3
 - The fourth order statistic is 5

Solution

Lemma

Let T be a binary tree with root r and n nodes, K a set of n distinct keys, ℓ the left child of r and s_ℓ the number of nodes in the subtree rooted at ℓ . If an assignment of labels from K to the nodes of T satisfies the binary search tree property, then $\text{label}(r)$ is the $s_\ell + 1$ order statistic of K .

Proof.

- For every node v of T in the left subtree of r it holds that $\text{label}(v) < \text{label}(r)$
 - T maintain the binary search tree property
- For every node v of T in the right subtree of r it holds $\text{label}(v) > \text{label}(r)$
 - T maintain the binary search tree property
- In the right subtree of r there are $n - s_\ell - 1$ nodes
- $\text{label}(r)$ is larger then s_ℓ keys in K and smaller then $n - s_\ell - 1$
- $\text{label}(r)$ is the $s_\ell + 1$ order statistic of K

Solution — cont.

Claim: There exists a single mapping of keys from K to nodes of T such that the resulting labeled binary tree maintains the binary search tree property

Proof. By induction on h , the height of T

- **Base:** holds for $h = 0$ since there is only one node and key
- **Hypothesis:** The claim holds for binary tree T with height $< h$
- **Step:**
- Let r be the root of T , ℓ its left child and i its right child
- Let s_ℓ be the number of nodes in the subtree rooted at ℓ
- Let x be the $s_\ell + 1$ order statistic of K
- By Lemma $\text{label}(r) = x$
- Let $L = \{k \in K \mid k < x\}$ and $I = \{k \in K \mid k > x\}$
 - $|L| = s_\ell$, $|I| = n - s_\ell - 1$

Solution — cont.

- The keys in the left subtree must be from L
- The keys in the right subtree must be from I
- The subtree rooted at ℓ has height $< h$. By induction hypothesis there exists a single mapping of keys from L to the nodes of the subtree rooted at ℓ such that the resulting labeled tree satisfies the binary search tree property
- The subtree rooted at i has height $< h$. By induction hypothesis there exists a single mapping of keys from I to the nodes of the subtree rooted at i such that the resulting labeled tree satisfies the binary search tree property

1 Binary search trees

2 Augmenting Data Structures

3 Non Distinct Keys

Rank and Select Operations

- Given a set S of n elements
- The $\text{rank}(x)$ operation on an element $x \in S$ returns its **position** in the linear order of the elements
- The $\text{select}(i)$ operation returns the i th order statistic of S

Example:

- $S = \{10, 1, 5, 18, 7, 3, 4\}$. The linear order is $(1, 3, 4, 5, 7, 10, 18)$
- $\text{rank}(1) = 1$, $\text{rank}(3) = 2$, $\text{rank}(18) = 7$
- $\text{select}(1) = 1$, $\text{select}(2) = 3$, $\text{select}(7) = 18$

Question 2

Design a dynamic data structure D with distinct keys that supports operations

- $\text{Init}(D)$ – initialize a new and empty data structure
- $\text{Insert}(D, x)$ – insert a new object x into D
- $\text{Delete}(D, x)$ – delete object x from D
- $\text{Rank}(x)$ – returns the position of $x.\text{key}$ in the linear order of keys in D

such that the run time of $\text{Init}(D)$ is $O(1)$ and the run time of $\text{Insert}(D, x)$, $\text{Delete}(D, x)$, and $\text{Rank}(x)$ is $O(\log n)$ where n is the number of elements currently stored in the data structure.

Augmenting Data Structures

- Some new required operations on a data set may be implemented using a "textbook" data structure
 - Doubly linked list, queue, binary search tree, 2-3 tree,...
- Usually no need to create an entirely new type of data structure
- Often it will suffice to augment a textbook data structure by storing additional information in it
- Augmenting a data structure is not always straightforward
 - Added information must be updated and maintained
- Suggested process:
 - ① Choose an underlying data structure
 - ② Determine additional information to maintain in the underlying data structure
 - ③ Verify that we can maintain the additional information for the basic modifying operations on the underlying data structure
 - ④ Develop new operations

2-3 tree

2-3 tree (2-3 יער)

- Rooted tree T satisfying:
 - Each internal node has **degree 2 or 3**
 - All leaves are at the **same level**
- Objects are stored **only** at the leaves
 - Internal nodes are there for organization purposes
- Data structure's attribute:
 - **root** = pointer to root of T
- Node attributes:
 - **left** = pointer to left child in T
 - **middle** = pointer to middle child in T
 - **right** = pointer to right child in T
 - **p** = pointer to parent in T
 - **key** = maximum key in its subtree
 - Leaves have additional satellite attributes (storing the actual data)
- When degree is 2, **left** and **middle** children are in use
- The **2-3 tree property**:
keys in left subtree < keys in middle subtree < keys in right subtree

Solution

- We will show an implementation based on a 2-3 tree
 - D will be an augmented 2-3 tree
- Each node v of the 2-3 tree will have (in addition) an attribute $size$
 - $v.size = \#$ of leaves (other than the sentinels) in the subtree rooted at v (including v)
 - Updating the value of $size$ attribute will occur during $\text{Insert}(D, x)$ and $\text{Delete}(D, x)$ operations in $O(\log n)$ time each

$\text{Init}(D)$: (high-level)

- Similar to 2_3_Init with the addition that the $size$ attribute of sentinels and root is set to 0

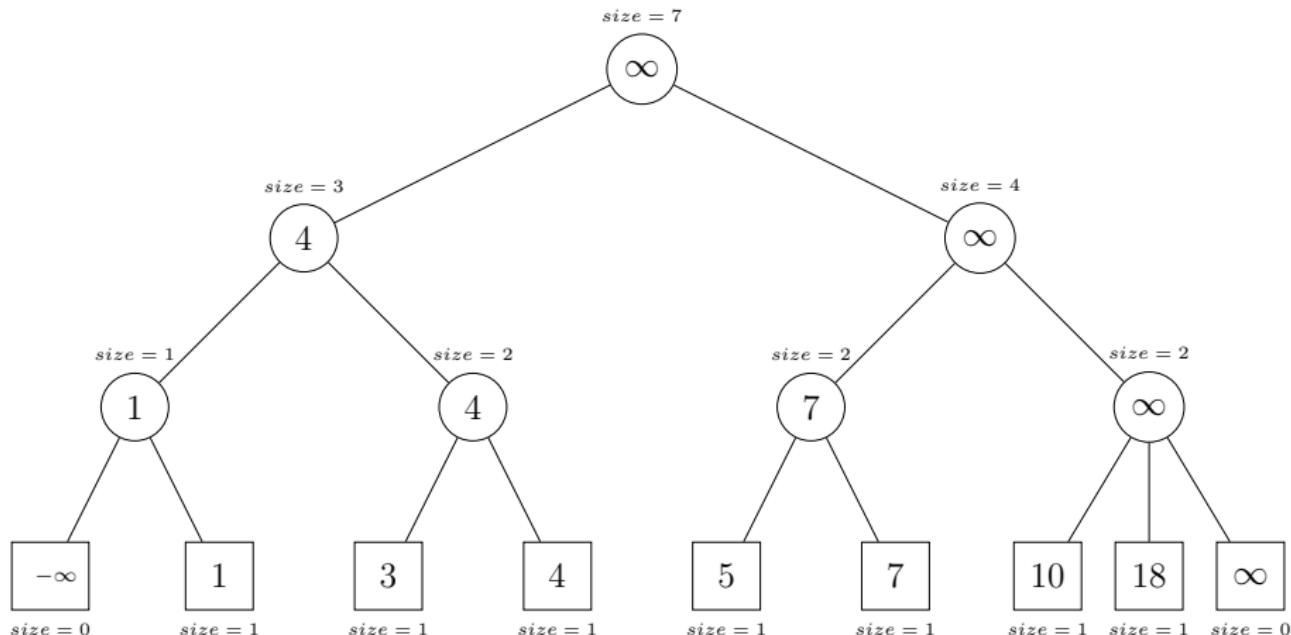
Solution — cont.

Returns the position of $x.key$ in the linear order of elements in D

Rank(x)

```
1: rank = 1
2: y = x.p
3: while y ≠ Nil do
4:     if x == y.middle then
5:         rank = rank + y.left.size
6:     else if x == y.right then
7:         rank = rank + y.left.size + y.middle.size
8:     x = y
9:     y=y.p
10:    return rank
```

Solution — example



Solution — cont.

Complexity Analysis:

- Each iteration of the while loop – one level up in the 2-3 tree
- In total – $O(\text{height}(D))$
- Recall, the height of a 2-3 tree is $O(\log n)$
- Time complexity of $\text{Rank}(x)$ is $O(\log n)$

Solution — cont.

Insert(D, x): (high-level)

- We will follow 2-3_Insert(T, z) to describe the insertion of a new node to D and the process of updating the *size* attributes
 - Recall, z is the node to be inserted
- Set $z.size = 1$
- 2-3_Insert(T, z) searches for the position to insert node z . Let y be its parent
- If $deg(y) = 2$, insert z and update $y.size = y.size + 1$ (in Insert_And_Split(x, z))
 - Update *size* values of all ancestors of y (add one to their previous value)
- Otherwise, update *size* attributes of y and the new node created during Insert_And_Split(x, z) (sum of children's *size* attribute)
 - Same problem in higher level
- If a new root needs to be added update its *size* value
 - Sum of children's *size* attribute
- We only added $O(1)$ time at each level in total $O(height(D))$

Solution — cont.

$\text{Delete}(D, x)$: (high-level)

- We will follow $\text{2_3_Delete}(T, z)$ to describe the deletion of a node from D and the process of updating the *size* attributes
 - Recall, z is the node to be deleted
- Let $y = z.p$ and delete z
- Update $y = y.size - 1$
- If $\text{deg}(y) = 2$ after deletion of z
 - Update *size* values of all ancestors of y (subtract one from their previous value)
- Otherwise, update *size* attribute of y and y sibling during $\text{Borrow_Or_Merge}(y)$ (sum of children's *size* attribute)
 - Same problem at higher level
- We only added $O(1)$ time at each level In total $O(\text{height}(D))$

Question 3

Design a dynamic data structure D with distinct keys that supports operations

- $\text{Init}(D)$ – initialize a new and empty data structure
- $\text{Insert}(D, x)$ – insert a new object x into D
- $\text{Delete}(D, x)$ – delete object x from D
- $\text{Select}(D, i)$ – returns a pointer to the element with the i th smallest key in D if exists, otherwise returns NIL

such that the run time of $\text{Init}(D)$ is $O(1)$ and the run time of $\text{Insert}(D, x)$, $\text{Delete}(D, x)$, and $\text{Select}(D, i)$ is $O(\log n)$ where n is the number of elements currently stored in the data structure.

Solution

- We will use 2-3 tree to implement the desired dynamic data structure with additional attribute *size* as defined in question 2
- Since we maintain the same attribute *size* we will use the same $\text{Insert}(D, x)$, $\text{Delete}(D, x)$ and, $\text{Init}(D)$ as in question 2

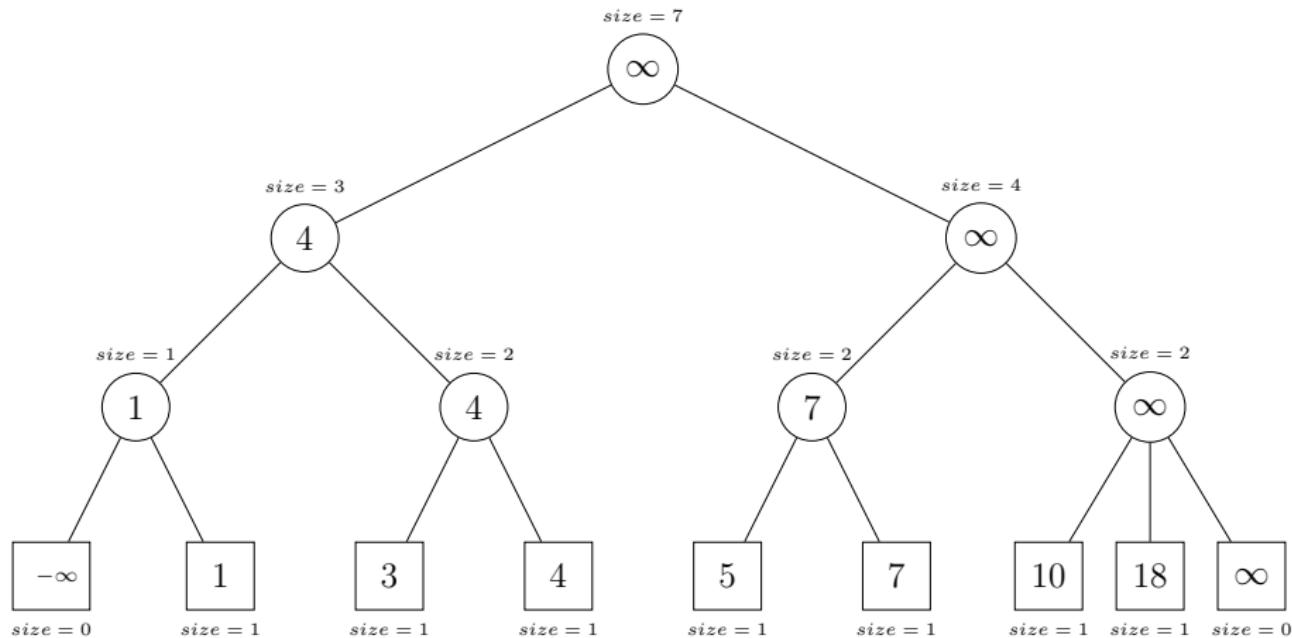
Solution — cont.

Returns a pointer to the element with the i th smallest key in the 2-3 tree rooted at x if such key exists, otherwise returns NIL

`Select_Rec(x, i)`

```
1: if  $x.\text{size} < i$  then
2:     return  $\text{NIL}$ 
3: if  $x$  is a leaf then
4:     return  $x$ 
5:  $s\_left = x.\text{left.size}$ 
6:  $s\_left\_middle = x.\text{middle.size} + x.\text{left.size}$ 
7: if  $i \leq s\_left$  then
8:     return Select_Rec( $x.\text{left}, i$ )
9: else if  $i \leq s\_left\_middle$  then
10:    return Select_Rec( $x.\text{middle}, i - s\_left$ )
11: else return Select_Rec( $x.\text{right}, i - s\_left\_middle$ )
```

Solution — example



Solution — cont.

Notice:

- In order to implement $\text{Select}(D, i)$ we need to run $\text{Select_Rec}(D.\text{root}, i)$

Complexity Analysis:

- Each recursion call – one level down in the 2-3 tree rooted at x
- In total – $O(\text{height}(x))$
- Recall, the height of 2-3 tree is $O(\log n)$
- Time complexity of $\text{Select}(D, i)$ is $O(\log n)$

Question 4

Design a dynamic data structure D with distinct keys that supports operations

- $\text{Init}(D)$ – initialize a new and empty data structure
- $\text{Insert}(D, x)$ – insert the element x into D
- $\text{Delete}(D, x)$ – delete element x from D
- $\text{Sum_of_smaller}(D, k)$ – returns the sum of keys in D that are less or equal to k

such that the run time of $\text{Init}(D)$ is $O(1)$ and the run time of $\text{Insert}(D, x)$, $\text{Delete}(D, x)$, and $\text{Sum_of_smaller}(D, k)$ is $O(\log n)$ where n is the number of elements currently stored in the data structure.

Solution

- We will show an implementation based on a 2-3 tree
 - D will be an augmented 2-3 tree
- Each node v of the 2-3 tree will have (in addition) an attribute sum
 - $v.sum =$ sum of keys in the leaves (other than the sentinels) in the subtree rooted at v (including v)
 - Updating the attribute sum will occur during $\text{Insert}(D, x)$ and $\text{Delete}(D, x)$ operations in $O(\log n)$ time each

$\text{Init}(D)$: (high-level)

- similar to 2_3_Init with the addition that the sum attribute of sentinels and root is set to 0

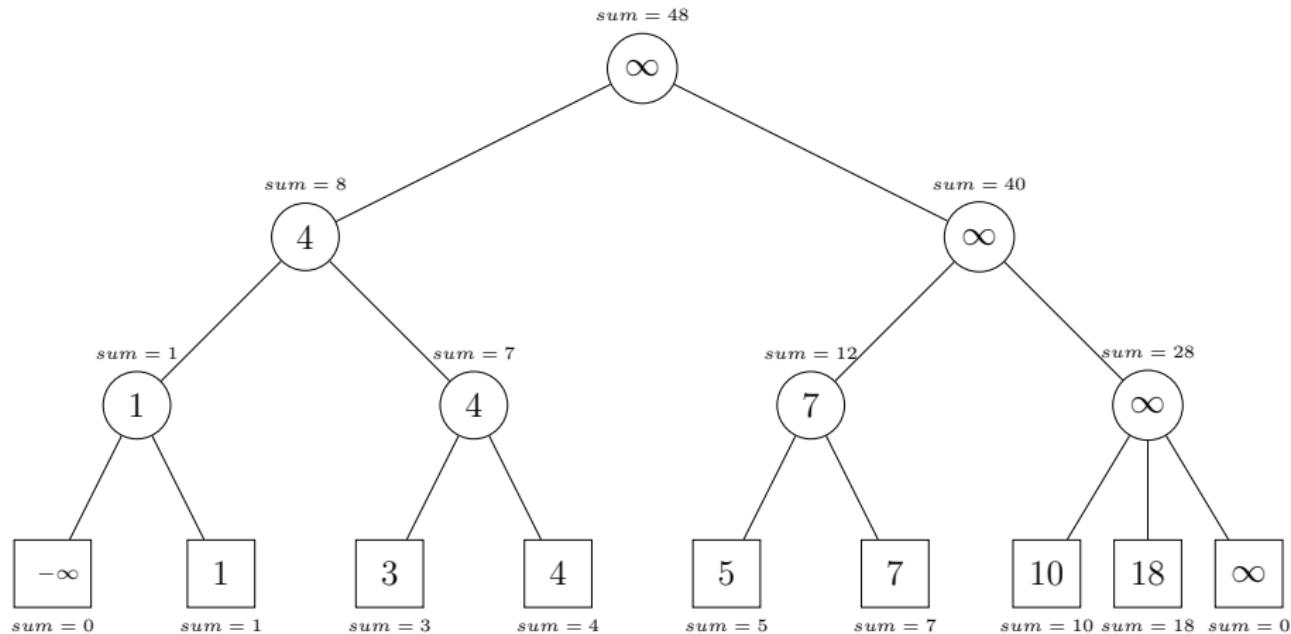
Solution — cont.

Returns the sum of keys that are $\leq k$ in the subtree rooted at x

`Sum_Of_Smaller_Rec(x, k)`

```
1: if  $x$  is a leaf then
2:   if  $x.key \neq -\infty$  then
3:     if  $x.key \leq k$  then
4:       return  $x.key$ 
5:     else
6:       return 0
7:   else
8:     return: Error no keys smaller than  $k$ 
9: if  $k \leq x.left.key$  then
10:   return Sum_Of_Smaller_Rec( $x.left, k$ )
11: else if  $k \leq x.middle.key$  then
12:   return  $x.left.sum + Sum\_Of\_Smaller\_Rec(x.middle, k)$ 
13: else return
       $x.left.sum + x.middle.sum + Sum\_Of\_Smaller\_Rec(x.right, k)$ 
```

Solution — example



Solution — cont.

Notice:

- In order to implement `Sum_Of_Smaller(D, k)` we need to run `Sum_Of_Smaller_Rec(D.root, k)`

Complexity Analysis:

- Each recursion call – one level down in the 2-3 tree rooted at x
- In total – $O(\text{height}(x))$
- Recall, the height of 2-3 tree is $O(\log n)$
- Time complexity of `2_3_Select(x, i)` is $O(\log n)$

Solution — cont.

$\text{Insert}(D, x)$: (high-level)

- We will follow $\text{2_3_Insert}(T, z)$ to describe the insertion of a new node to D and the process of updating the sum attributes
 - Recall, z is the node to be inserted
- Set $z.sum = z.key$
- $\text{2_3_Insert}(T, z)$ searches for the position to insert node z . Let y be its parent
- If $\text{deg}(y) = 2$, insert z and update $y.sum = y.sum + z.key$ (in $\text{Insert_And_Split}(x, z)$)
 - Update sum values of all ancestors of y (add $z.key$ to their previous value)
- Otherwise, update sum attributes of y and the new node created during $\text{Insert_And_Split}(x, z)$ (sum of children's sum attribute)
 - Same problem in higher level
- If a new root needs to be added update its sum value
 - Sum of children's sum attribute
- We only added $O(1)$ time at each level In total $O(\text{height}(D))$

Solution — cont.

$\text{Delete}(D, x)$: (high-level)

- We will follow $\text{2_3_Delete}(T, z)$ to describe the deletion of a node from D and the process of updating the sum attributes
 - Recall, z is the node to be deleted
- Let $y = z.p$ and delete z
- Update $y = y.sum - z.key$
- If $\text{deg}(y) = 2$ after deletion of z
 - Update sum values of all ancestors of y (subtract $z.key$ from their previous value)
- Otherwise, update sum attribute of y and y 's sibling during $\text{Borrow_Or_Merge}(y)$ (sum of children's sum attribute)
 - Same problem at higher level
- We only added $O(1)$ time at each level In total $O(\text{height}(D))$

Question 5

Design a dynamic data structure D with distinct keys where each element has a numerical attribute *value* that supports operations

- $\text{Init}(D)$ – initialize a new and empty data structure
- $\text{Insert}(D, x)$ – insert the element x into D
- $\text{Delete}(D, x)$ – delete element x from D
- $\text{Largest_Value}(D)$ – returns the largest value of elements in D

such that the run time of $\text{Init}(D)$ and $\text{Largest_Value}(D)$ is $O(1)$ and the run time of $\text{Insert}(D, x)$ and $\text{Delete}(D, x)$ is $O(\log n)$ where n is the number of elements currently stored in the data structure.

Solution

- We will show an implementation based on a 2-3 tree
 - D will be an augmented 2-3 tree
- Each node v of the 2-3 tree will hold the attribute $value$
 - $v.value = \text{largest value in the subtree rooted at } v \text{ (including } v\text{)}$
 - Updating the attribute $value$ will occur during $\text{Insert}(D, x)$ and $\text{Delete}(D, x)$ operations in $O(\log n)$ time each

$\text{Init}(D)$: (high-level)

- Similar to 2_3_Init with the addition that the $value$ attribute of sentinels and root is set to $-\infty$

Solution — cont.

Returns the largest value in D

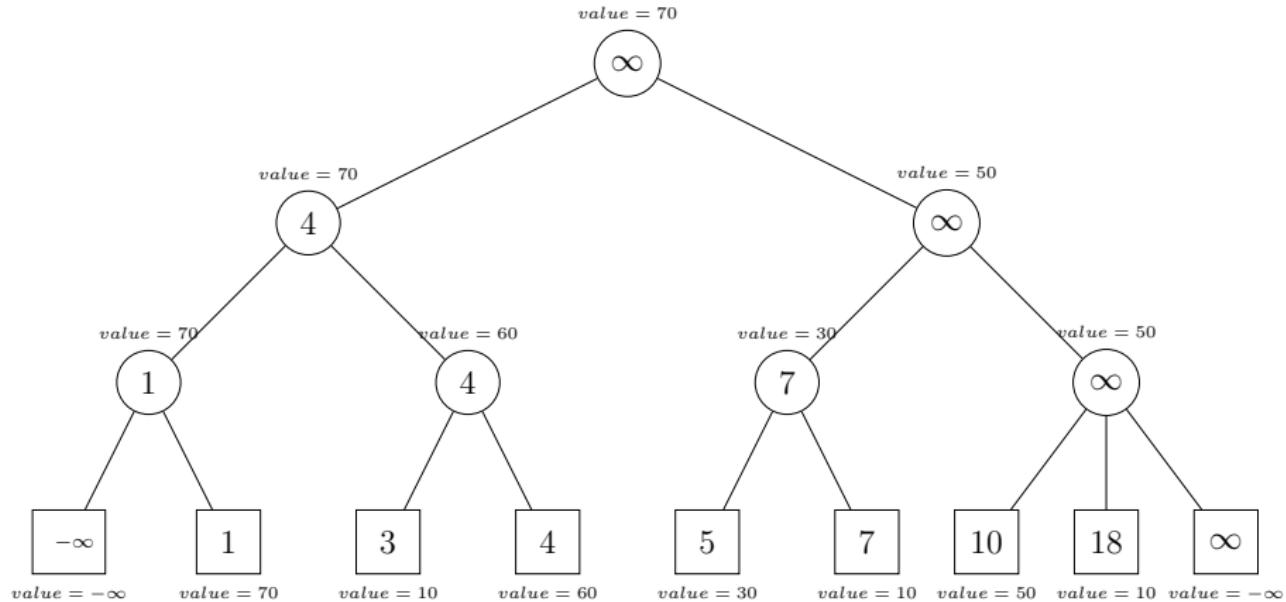
Largest_Value(D)

```
1: if  $D.root.value \neq -\infty$  then           ▷ check if  $D$  is empty
2:     return  $D.root.value$ 
3: else
4:     return error:  $D$  is empty
```

Complexity Analysis:

- Only constant time operations – $O(1)$

Solution — example



Solution — cont.

Insert(D, x): (high-level)

- We will follow `2_3_Insert(T, z)` to describe the insertion of a new node to D and the process of updating the *value* attributes
 - Recall, z is the node to be inserted
- `2_3_Insert(T, z)` searches for the position to insert node z . Let y be its parent
- If $\text{deg}(y) = 2$, insert z and update $y.\text{value} = \max_{w \in \text{Child}(y)}\{w.\text{value}\}$ (in `Insert_And_Split(x, z)`)
 - Update *value* of all ancestors of y (max of children's *value* attribute)
- Otherwise, update *value* attributes of y and the new node created during `Insert_And_Split(x, z)` (max of children's *value* attribute)
 - Same problem in higher level
- If a new root needs to be added update its *value* attribute
 - Max of children's *value* attribute
- We only added $O(1)$ time at each level In total $O(\text{height}(D))$

Solution — cont.

Delete(D, x): (high-level)

- We will follow 2_3_Delete(T, z) to describe the deletion of a node from D and the process of updating the *value* attributes
 - Recall, z is the node to be deleted
- Let $y = z.p$ and delete z
- Update $y.value = \max_{w \in Child(y)}\{w.value\}$
- If $\deg(y) = 2$ after z delete
 - Update *value* of all ancestors of y (max of children's *value*)
- Otherwise, update *value* attribute of y and y 's sibling during Borrow_Or_Merge(y) (max of children's *value* attribute)
 - Same problem at higher level
- We only added $O(1)$ time at each level In total $O(\text{height}(D))$

1 Binary search trees

2 Augmenting Data Structures

3 Non Distinct Keys

Non distinct keys

- All data structures we saw assumed uniqueness of keys
- This assumption can be easily lifted with slight modifications in the data structure operations
- This assumption can also be lifted without changing the data structure operations
 - For example:
 - We maintain an additional data structure attribute *count* initialized to 0 once the data structure is initialized
 - After each insertion of an object the count attribute is incremented by 1
 - Let k be some object's key to be inserted
 - We define the ordered pair (k, count) as the new key
 - Let (k_1, c_1) and (k_2, c_2) be two new keys
 - We define an order on keys by
$$(k_1, c_1) < (k_2, c_2) \text{ iff } k_1 < k_2 \vee (k_1 = k_2 \wedge c_1 < c_2)$$